

A background image of an artist with dark hair and a beard, wearing a dark shirt, painting a large, colorful abstract mural. The artist is holding a paintbrush in his right hand and a red paint can in his left hand. The mural features vibrant orange, red, and blue colors. The text "EMPOWER: YOU" is overlaid on the image.

EMPOWER: YOU

# Java Persistence API: Best Practices

Carol McDonald  
Java Architect

**SUN TECH DAYS 2008–2009**  
A Worldwide Developer Conference



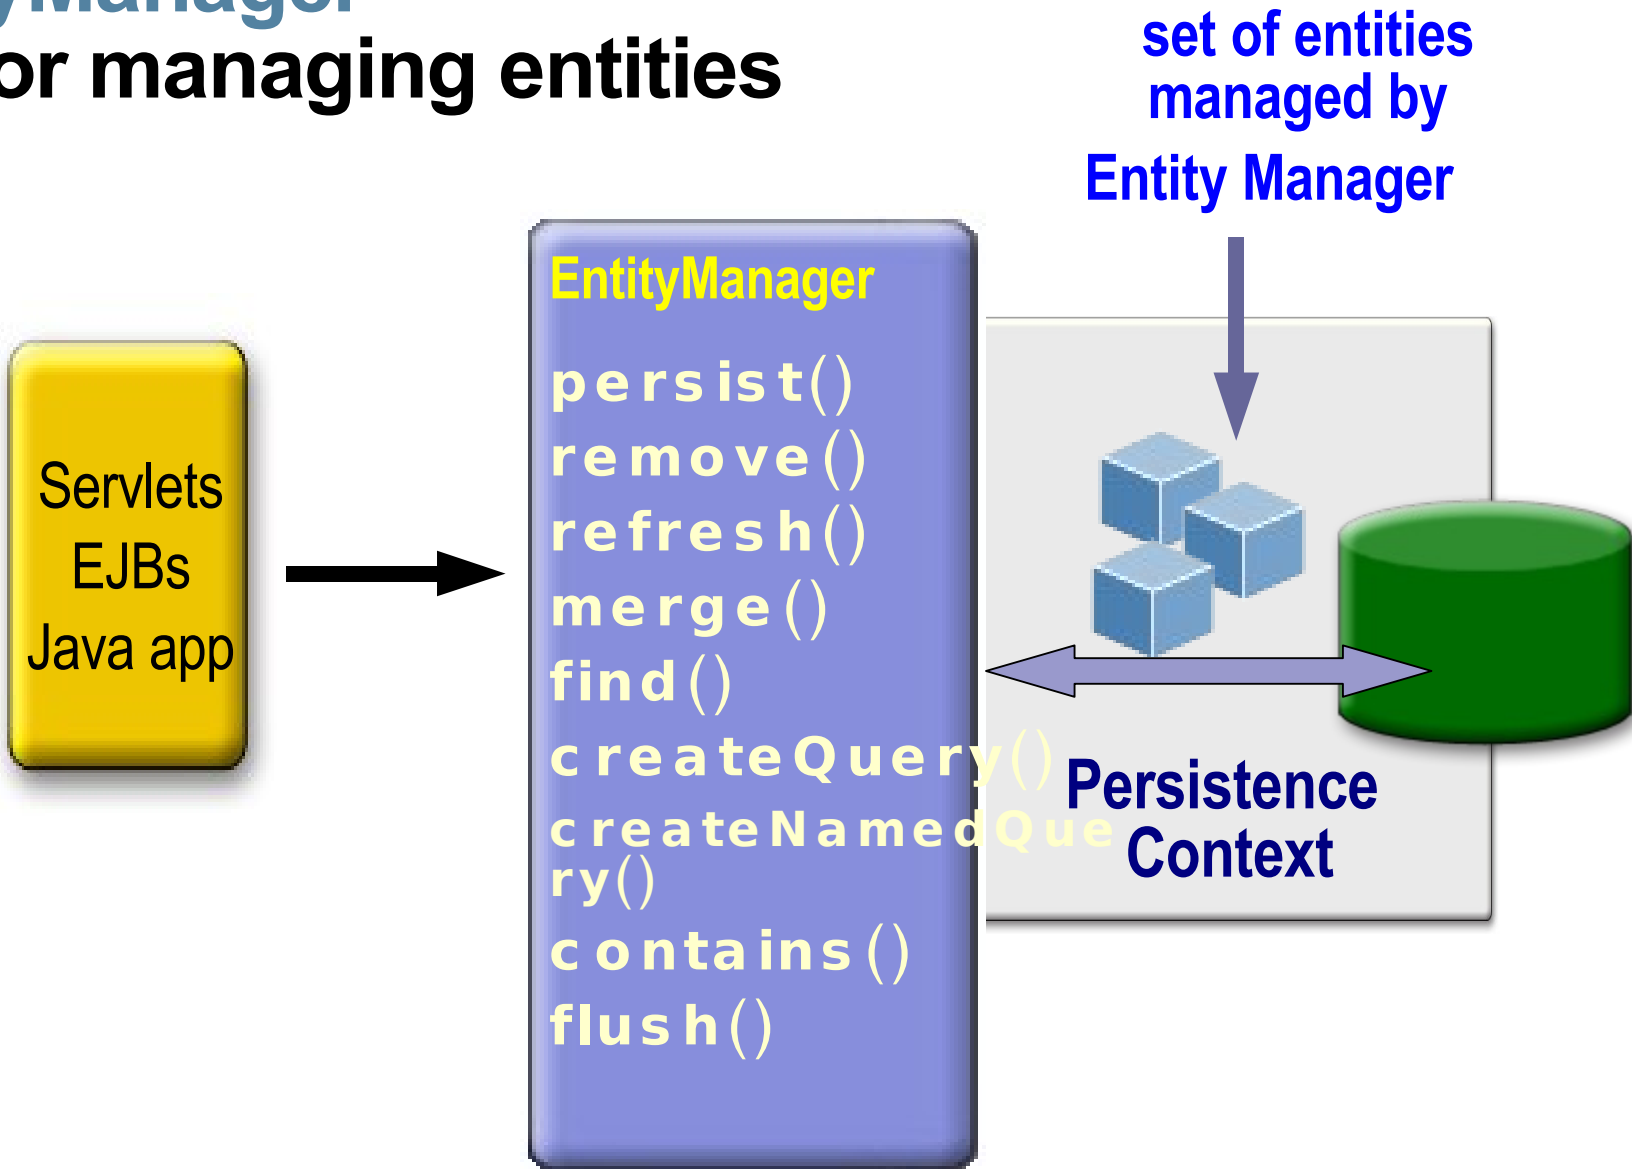
# Agenda

- > **Entity Manager**
- > Persistence Context
- > Entities
- > Schema & Queries
- > Transaction



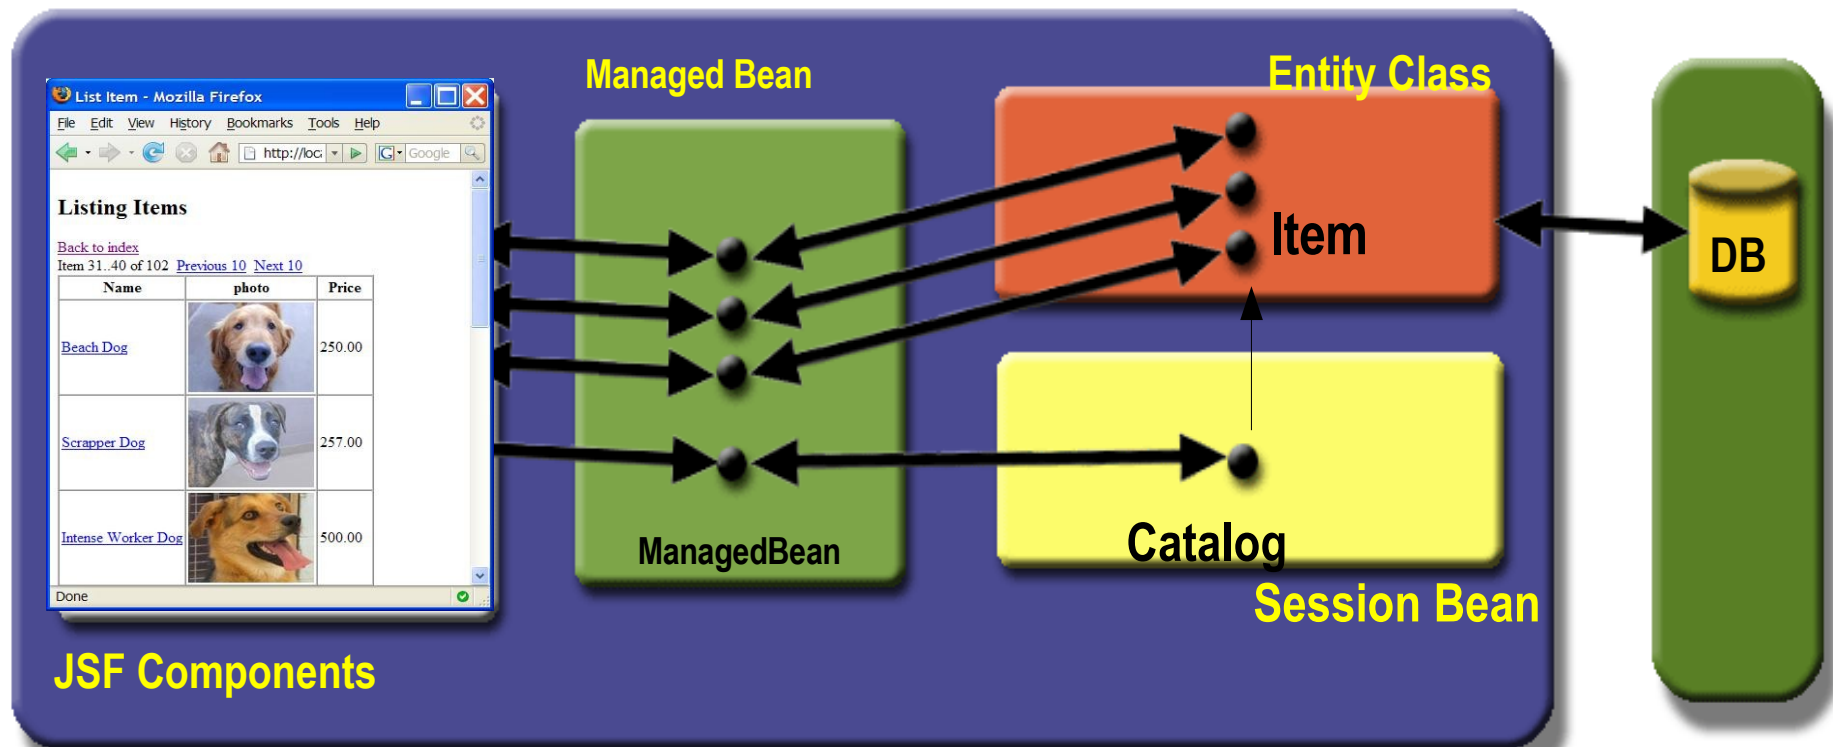
# EntityManager

## API for managing entities





# Catalog Java EE Application







# EJB EntityManager Example

Dependency Injection

**@Stateless**

public class **Catalog** implements **CatalogService** {

**@PersistenceContext(unitName="PetCatalogPu")**  
**EntityManager em;**



**@TransactionAttribute(NOT\_SUPPORTED)**

public List<Item> **getItems**(int **firstItem**,  
int **batchSize**) {

**Query q = em.createQuery**

("select i from Item as i");

**q.setMaxResults(batchSize);**

**q.setFirstResult(firstItem);**

**List<Item> items= q.getResultList();**

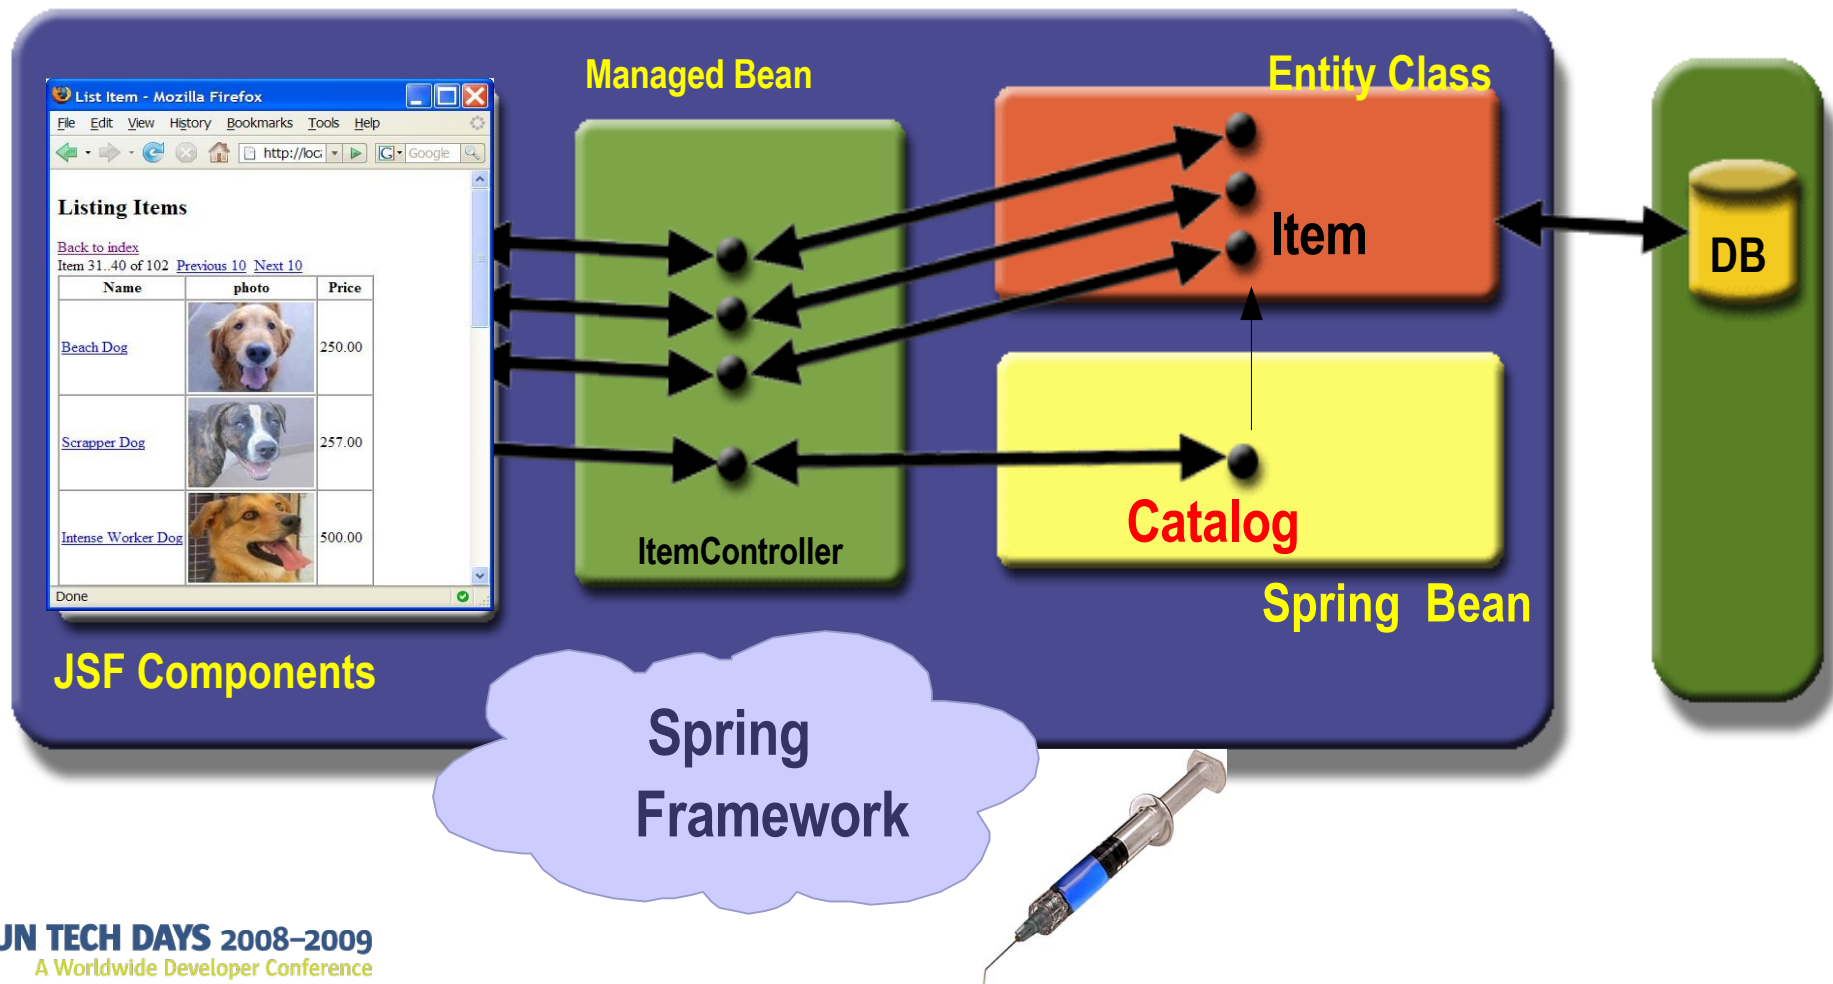
**return items;**

}

}



# Catalog Spring JPA Application






# Spring with JPA

Component Stereotype

Spring transactions use aop

```
@Repository
@Transactional
public class CatalogDAO implements CatalogService {
```

```
@PersistenceContext(unitName="PetCatalogPu")
private EntityManager em;
```



```
@Transactional(readOnly=true)
public List<Item>.getItems(int firstItem,int batchSize) {
    Query q =
        em.createQuery("select object(o) from Item as o");
    q.setMaxResults(batchSize);
    q.setFirstResult(firstItem);
    List<Item> items= q.getResultList();
    return items;
}
```



# Container vs Application Managed

Container managed entity managers (EJB, Spring Bean, Seam component)

- Injected into application
- Automatically closed
- JTA transaction – propagated

Application managed entity managers

- > Used outside of the JavaEE 5 platform
- > Need to be explicitly created
  - `Persistence.createEntityManagerFactory()`
- > `RESOURCE_LOCAL` transaction – not propagated
- > Need to explicitly close entity manager





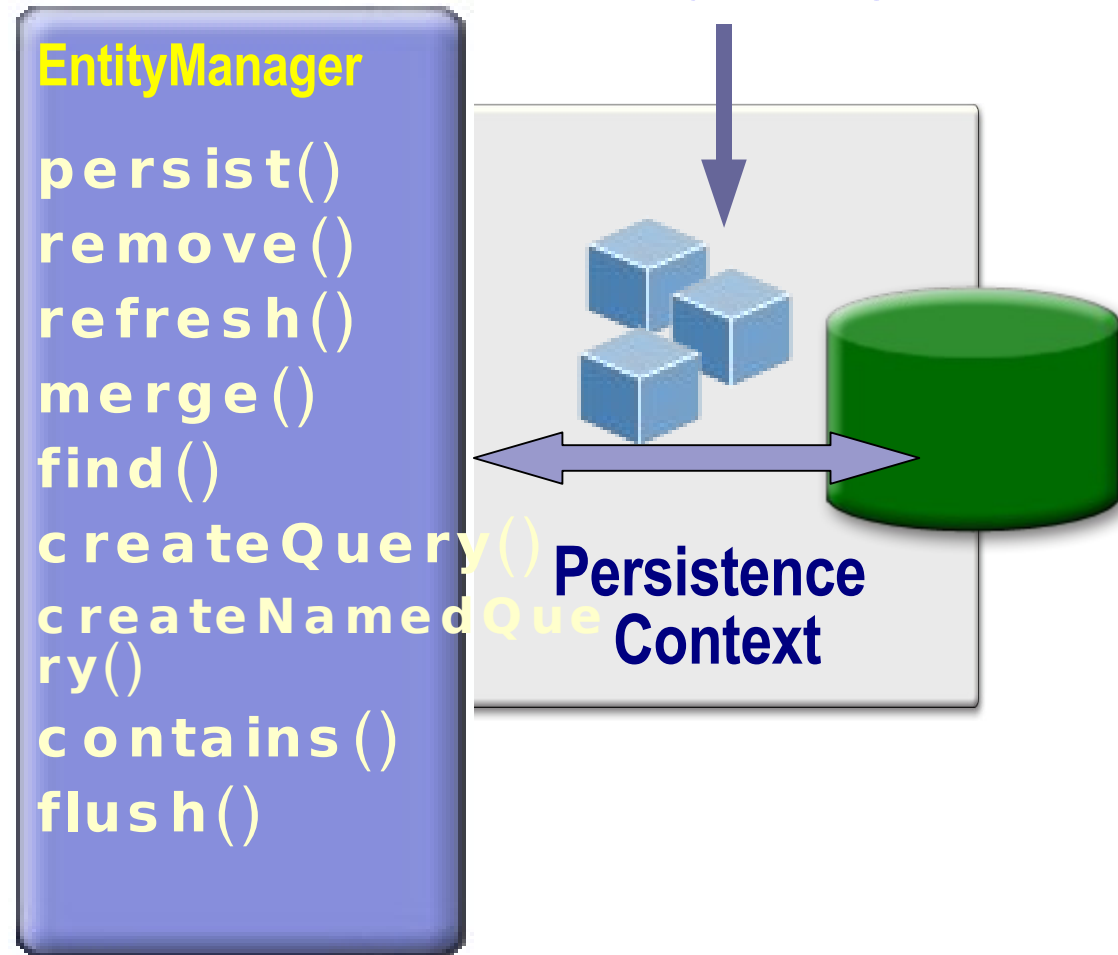
# Agenda

- > Entity Manager
- > **Persistence Context**
- > Entities
- > Queries
- > Transaction



# Persistence Context

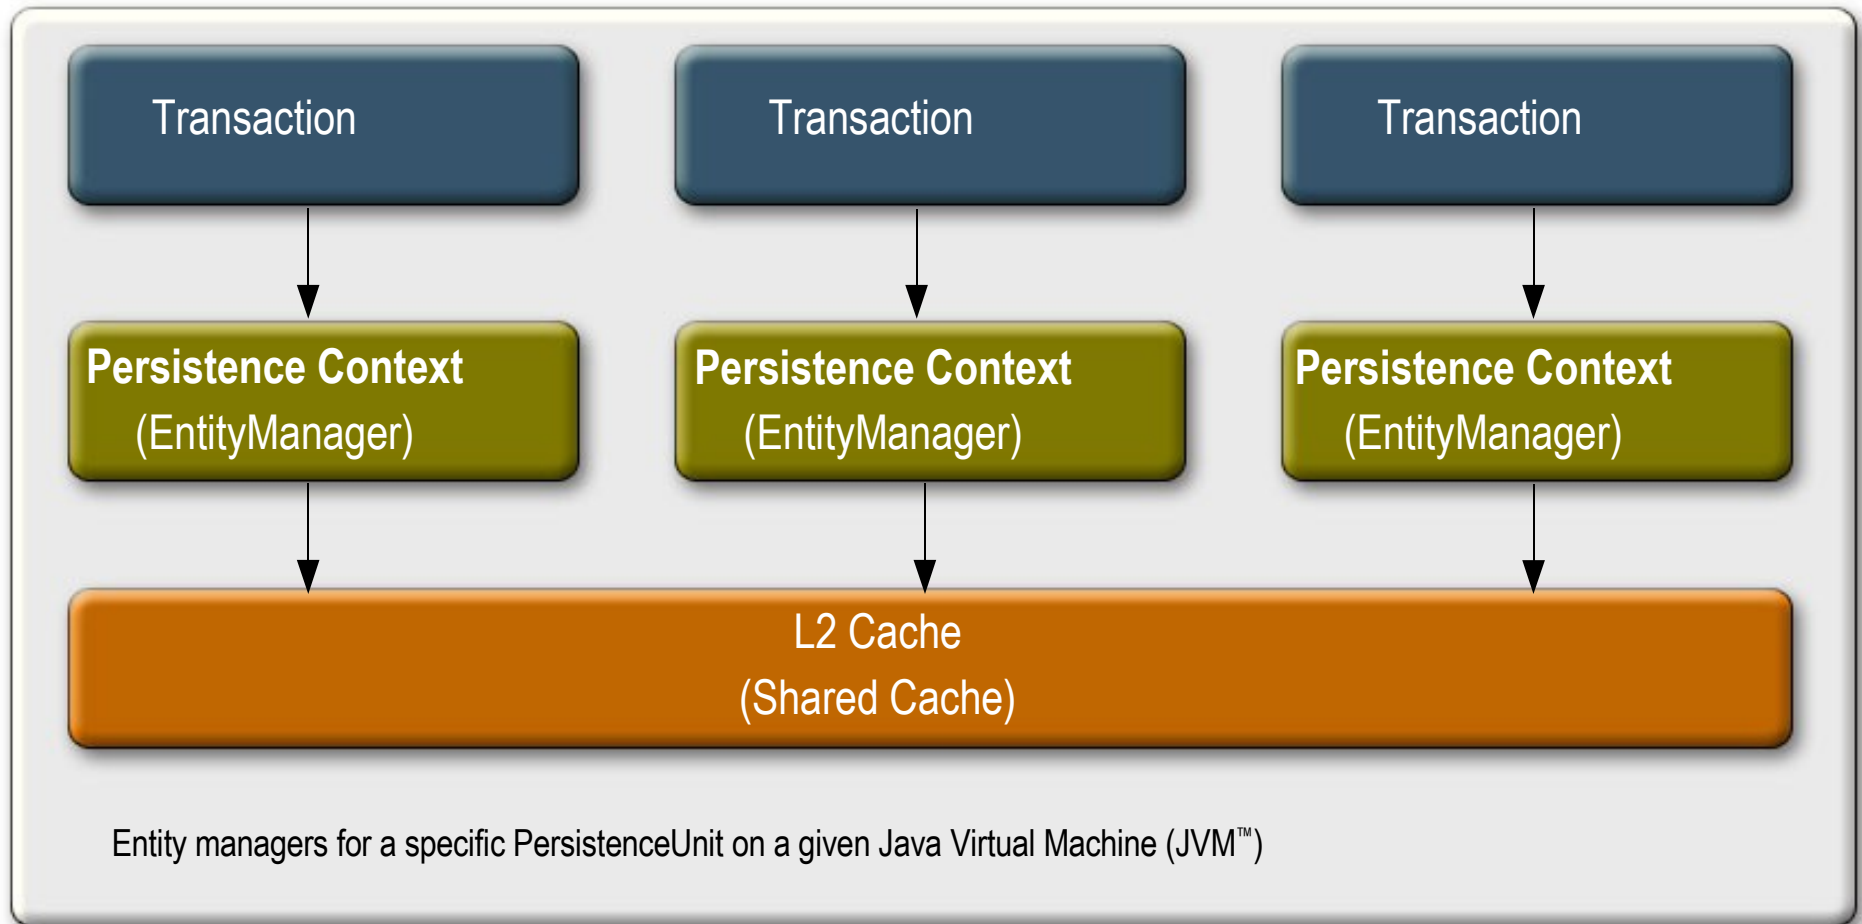
- Persistence context acts as a first level **cache** for entities
- Two types of persistence context
  - > Transaction scoped
  - > Extended scoped persistence context





# Level1 and Level2 caches

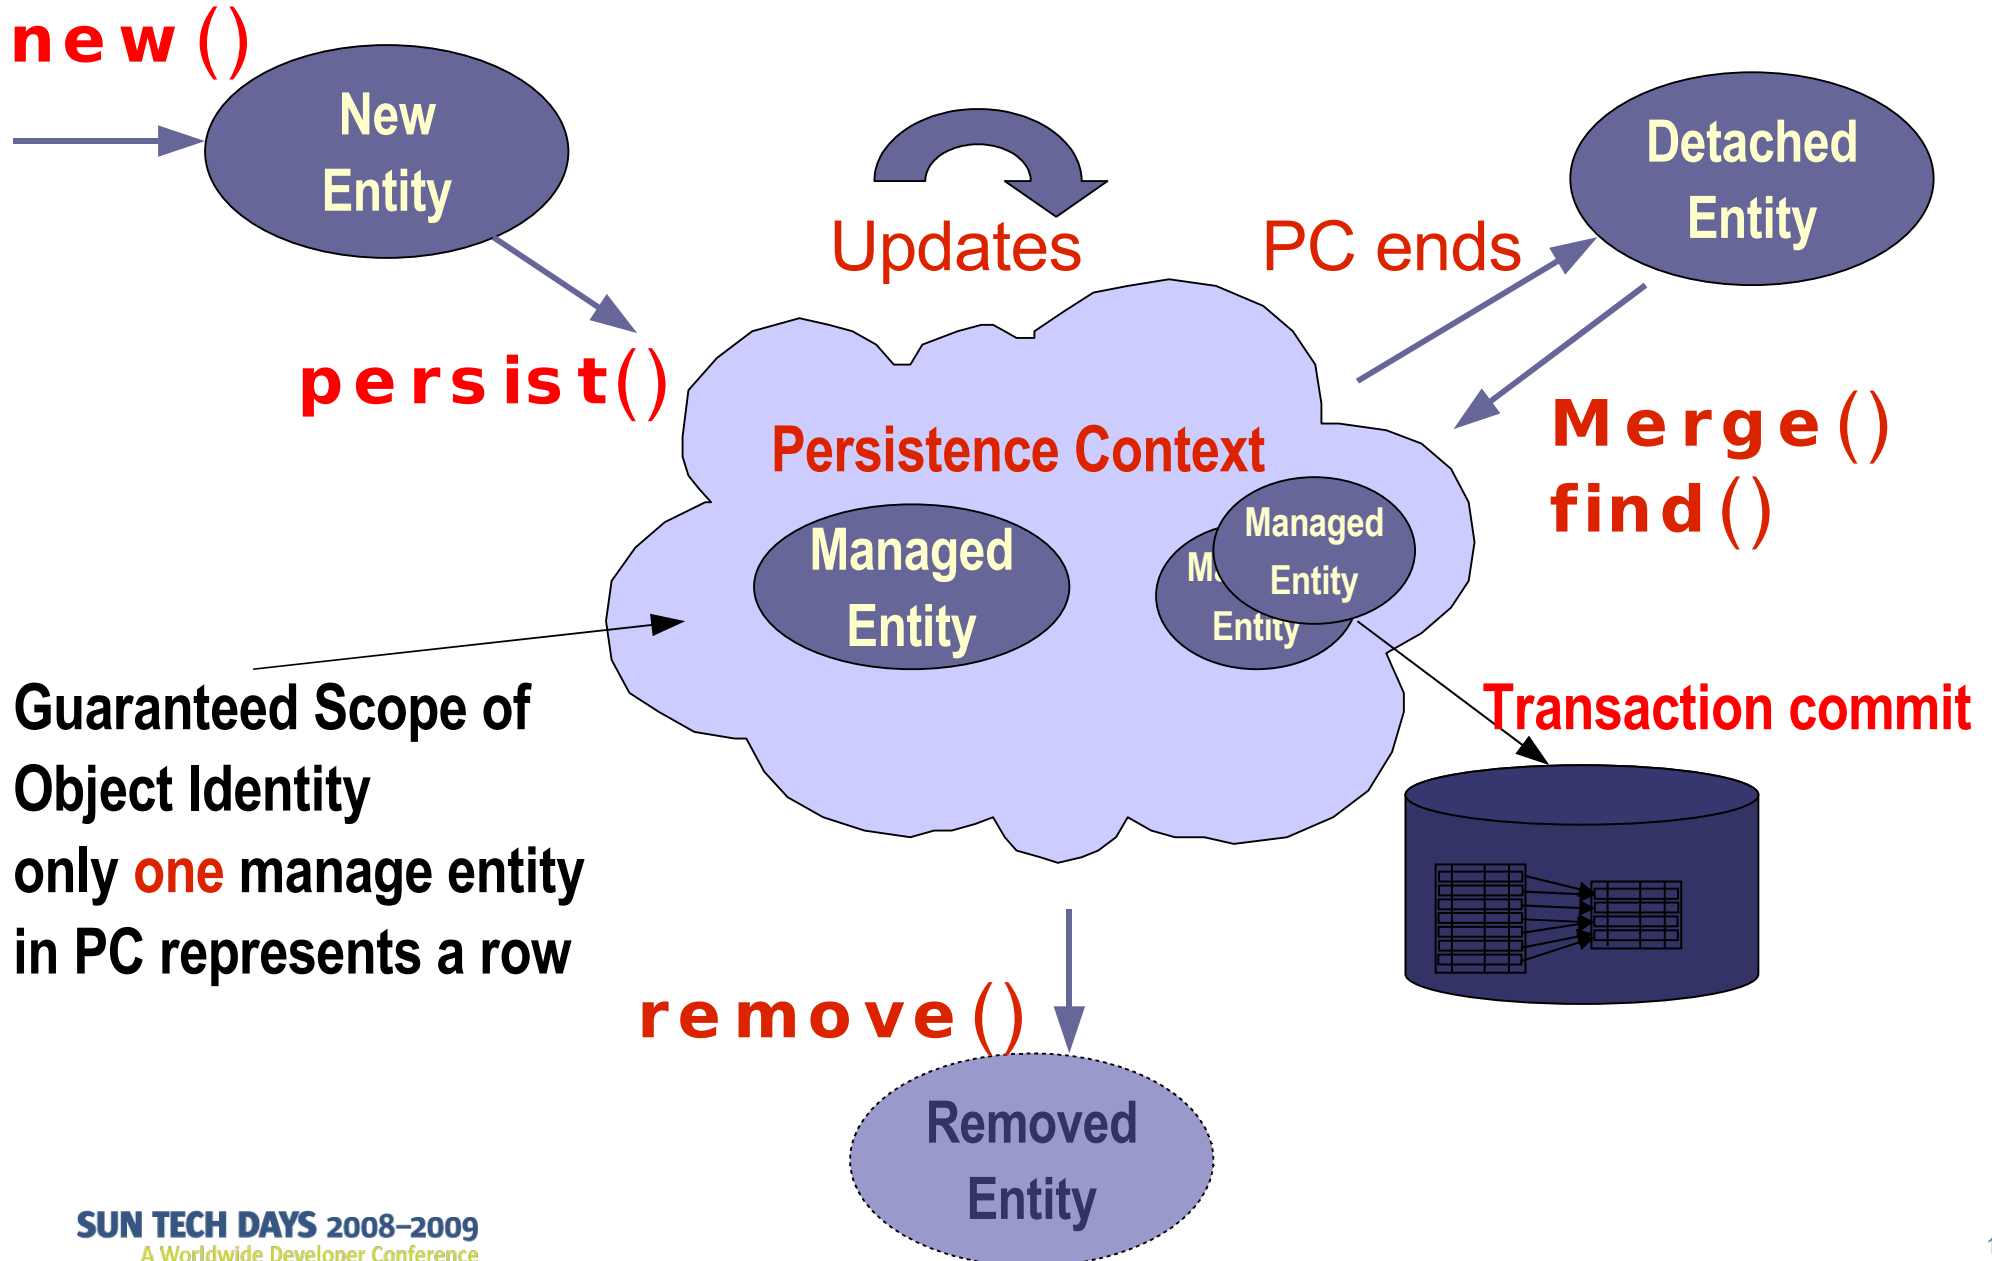
Persistence Context is a Level 1 cache



The terms “Java Virtual Machine” and “JVM” mean a Virtual Machine for the Java™ Platform.  
Source: [http://weblogs.java.net/blog/guruwons/archive/2006/09/understanding\\_t.html](http://weblogs.java.net/blog/guruwons/archive/2006/09/understanding_t.html)



# Entity Lifecycle





# Entity Lifecycle Illustrated – The Code

```

@Stateless public ShoppingCartBean
    implements ShoppingCart {

    @PersistenceContext EntityManager entityManager;

    public OrderLine createOrderLine(Product product
        , Order order) {
        OrderLine orderLine = new OrderLine(order, product);
        entityManager.persist(orderLine);
        return (orderLine);
    }
}

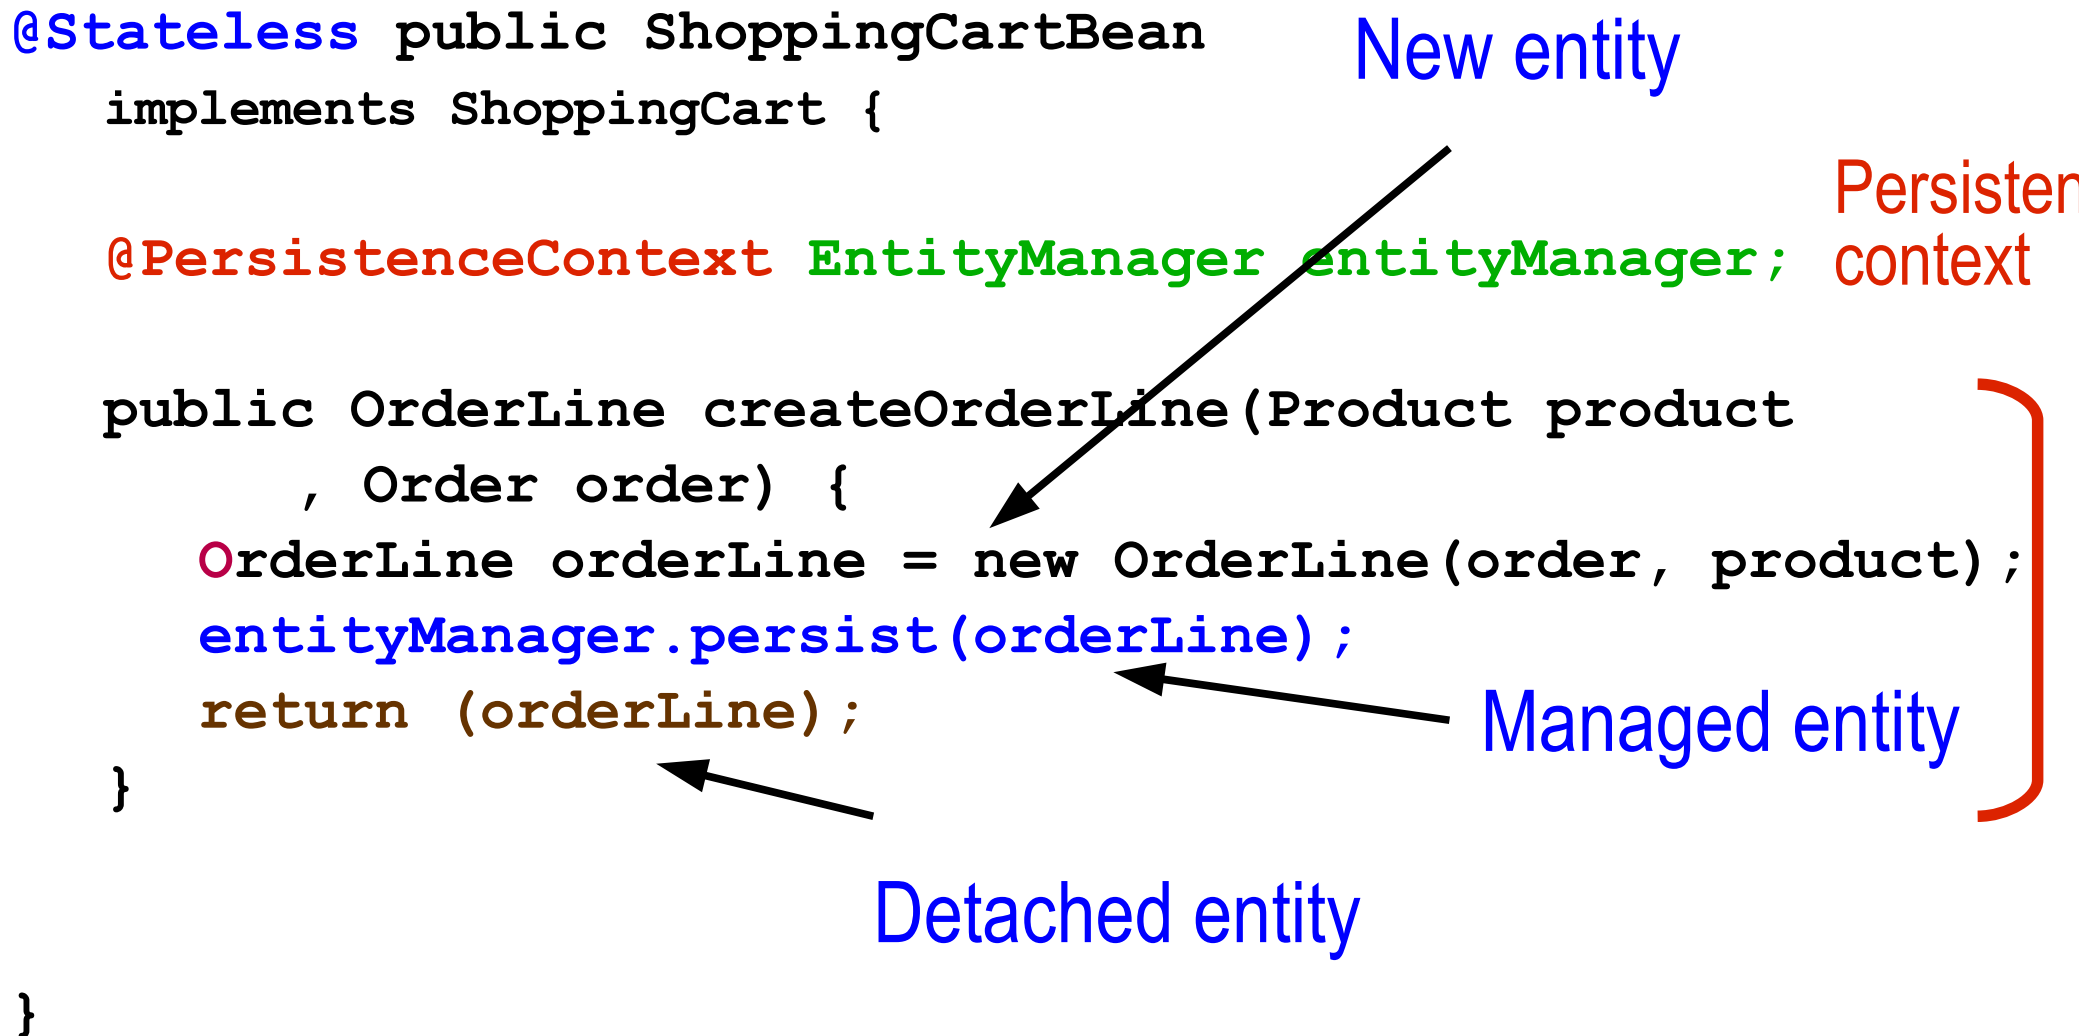
```

New entity

Persistence context

Managed entity

Detached entity







# Scope of Identity

```
@Stateless public ShoppingCartBean implements ShoppingCart {

    @PersistenceContext EntityManager entityManager; Persistence context

    public OrderLine createOrderLine(Product product, Order order) {
        OrderLine orderLine = new OrderLine(order, product);
        entityManager.persist(orderLine);
        OrderLine orderLine2 = entityManager.find(OrderLine,
            orderLine.getId());
        (orderLine == orderLine2) // TRUE
        return (orderLine);
    }
}
```

Multiple retrievals of the same object return references to the **same object instance**



# Persistence Context

- Two types of persistence context
- Transaction scoped
  - > Used in stateless components
  - > Typically begins/ends at request entry/exit points respectively
- Extended scoped persistence context



# Persistence Context Propagation

**@Stateless** public class **ShoppingCartBean** implements  
ShoppingCart {

**@EJB InventoryService inv;**

**@EJB OrderService ord;**

```
public void checkout(Item i, Product p) {
    inv.createOrder(item);
    ord.updateInventory(Product p)
```

```
}
```

```
}
```



Persistence context



# Persistence Context Propagation

```
@Stateless public class OrderServiceBean implements
    OrderService {
```

```
    @PersistenceContext EntityManager em1;
    public void createOrder(Item item) {
        em1.persist(new Order(item));
    }
}
```

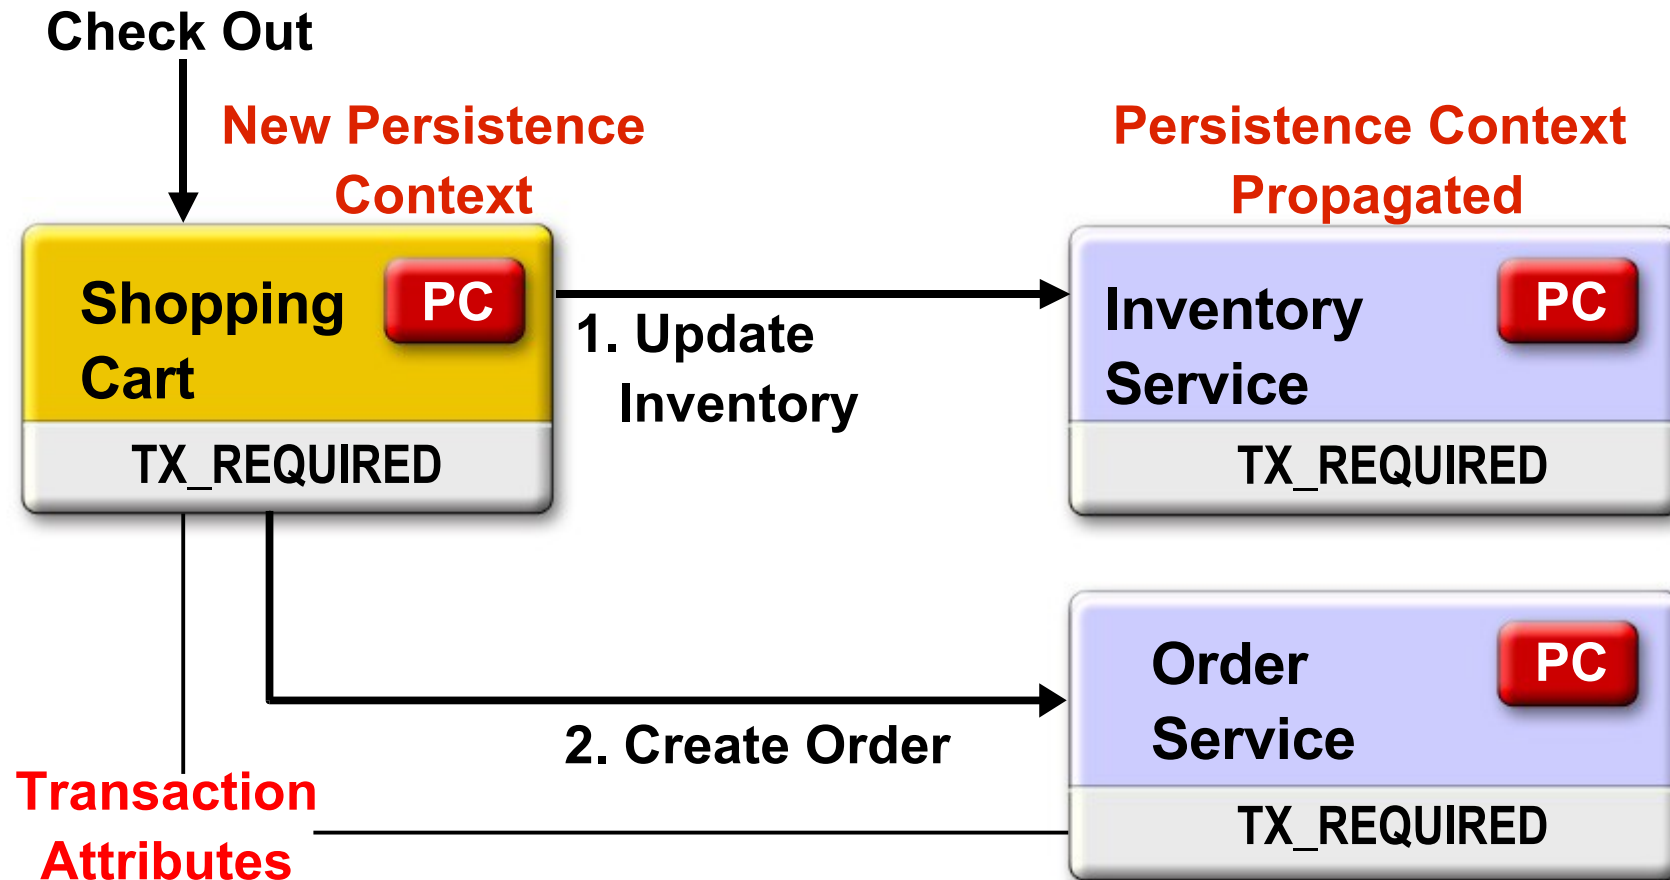
```
@Stateless public class InventoryServiceBean implements
    InventoryService {
```

```
    @PersistenceContext EntityManager em2;
    public void updateInventory(Product p) {
        Product product = em2.merge(p);
    }
}
```

...



# Declarative Transaction Management Example







# AuditServiceBean


**@Stateless**

**public class AuditServiceBean implements AuditService {**

**@PersistenceContext**

**private EntityManager em;**

**NEW  
PC !**



**@TransactionAttribute(REQUIRES\_NEW)**

**public void logTransaction2(int id, String action) {**

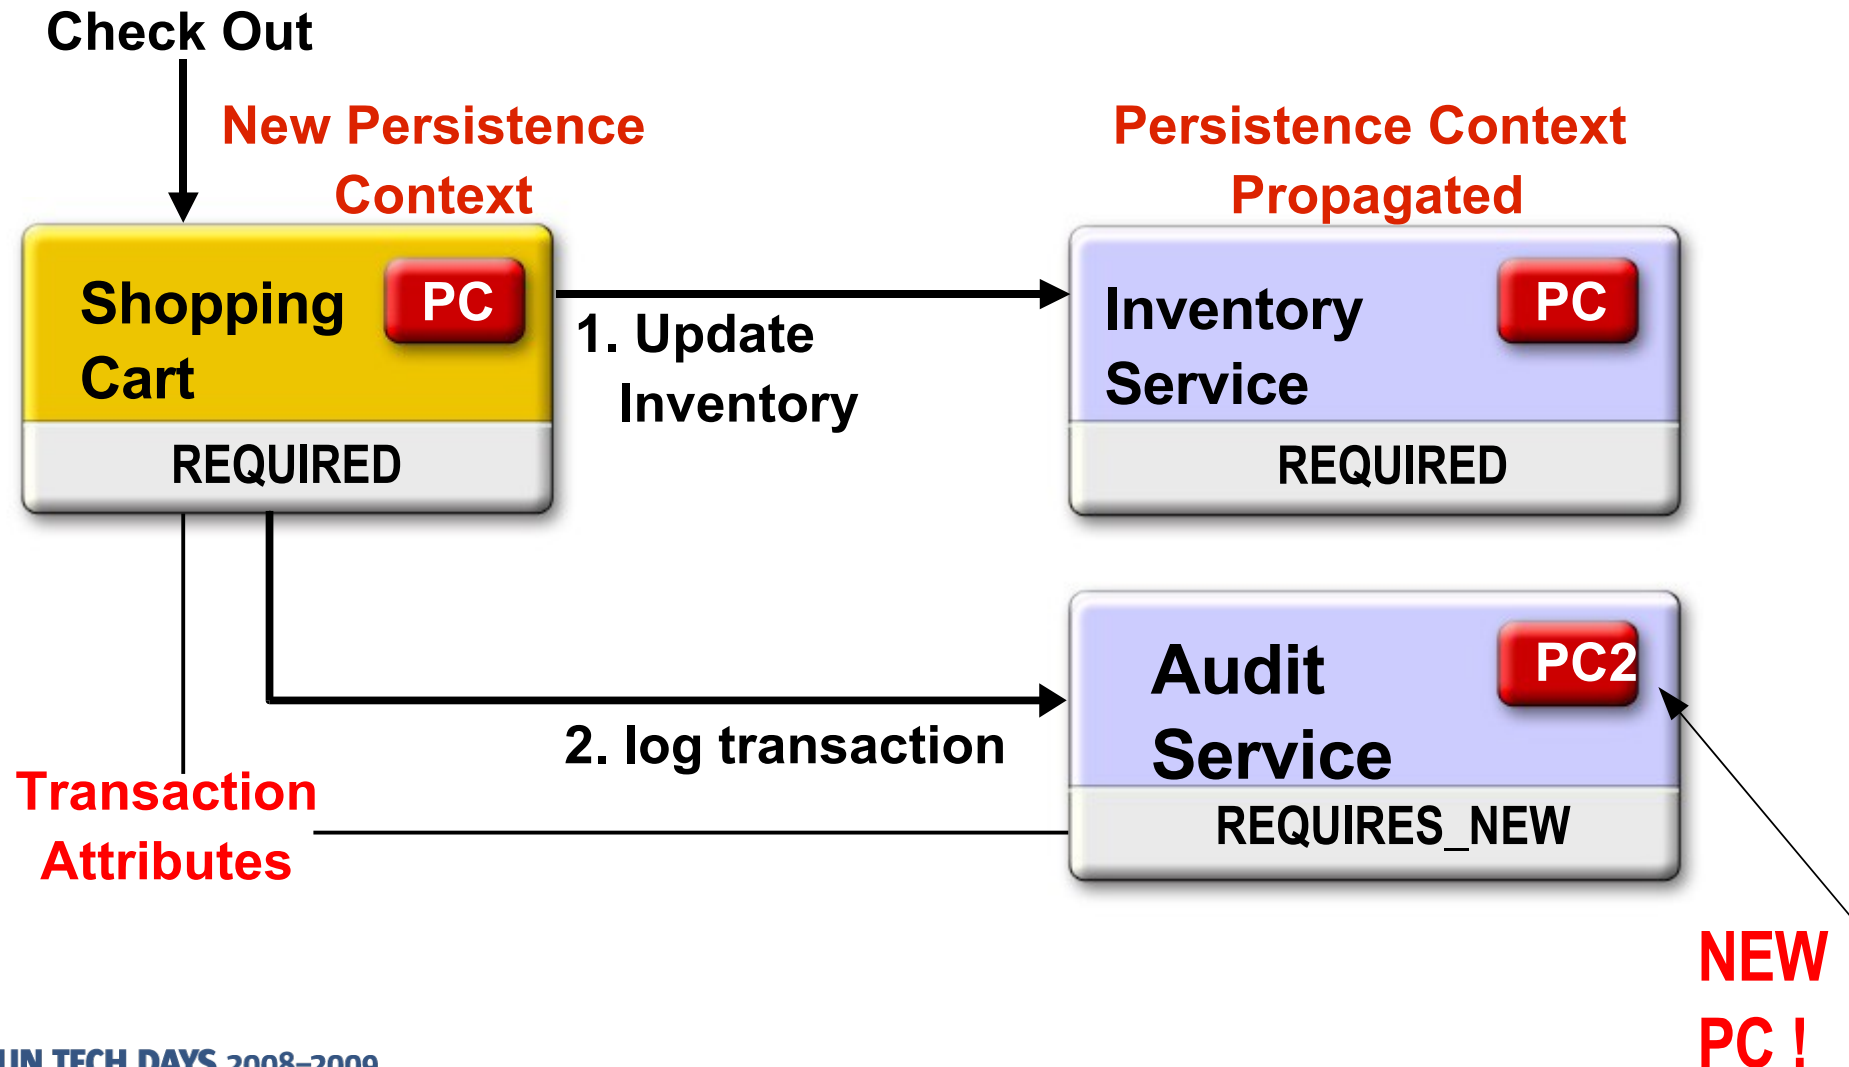
**LogRecord lr = new LogRecord(id, action);**

**em.persist(lr);**

**}**



# Declarative Transaction Management Example 2





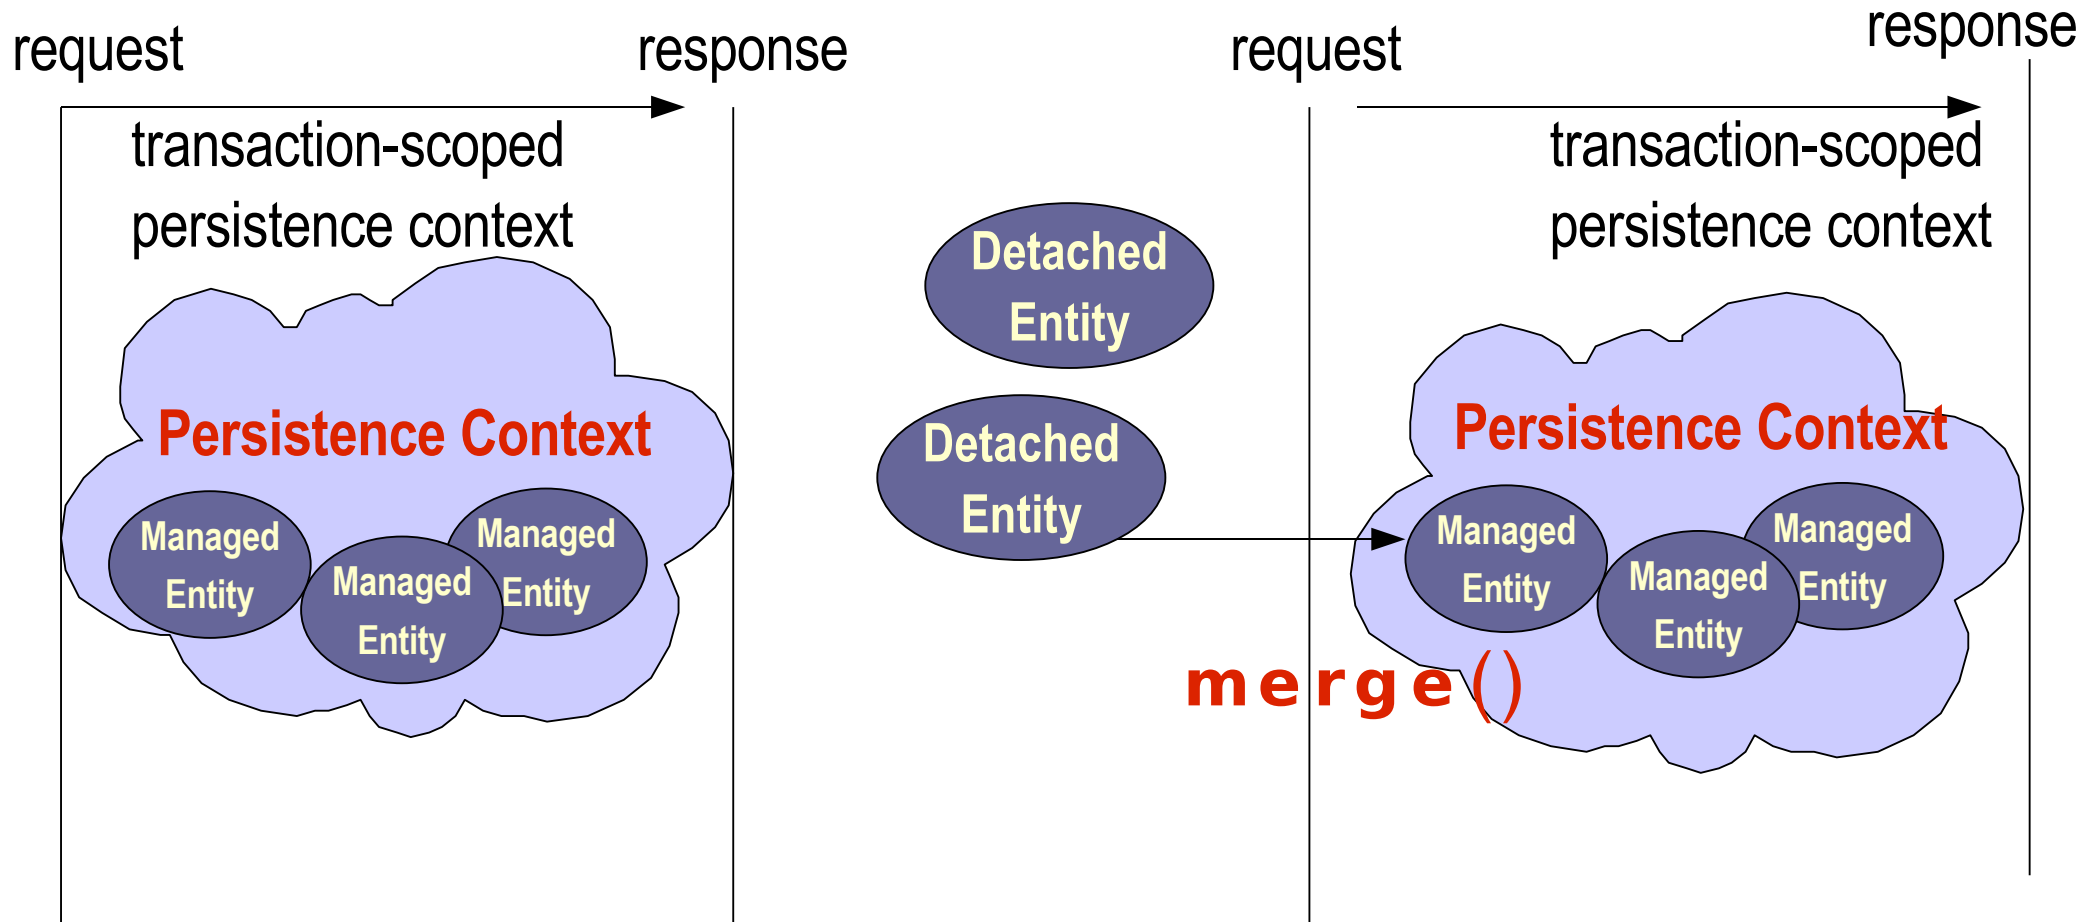
# Persistence Provider PC Transaction Features

- Attribute-level change tracking
- Only the minimal updates are sent to the database
- Orders INSERT, UPDATE and DELETE statements
- Minimizes database interactions
- EntityManager flush SQL prior to commit



# Conversation with detached entity

## Conversation





# Conversation with detached entity

```
@Stateless public ShoppingCartBean implements ShoppingCart {
    @PersistenceContext EntityManager entityManager;

    public OrderLine createOrderLine(Product product, Order order) {
        OrderLine orderLine = new OrderLine(order, product);
        entityManager.persist(orderLine);
        return (orderLine);
    }

    public OrderLine updateOrderLine(OrderLine orderLine) {
        OrderLine orderLine2 = entityManager.merge(orderLine);
        return orderLine2;
    }
}
```

Managed entity

Detached entity

Managed entity



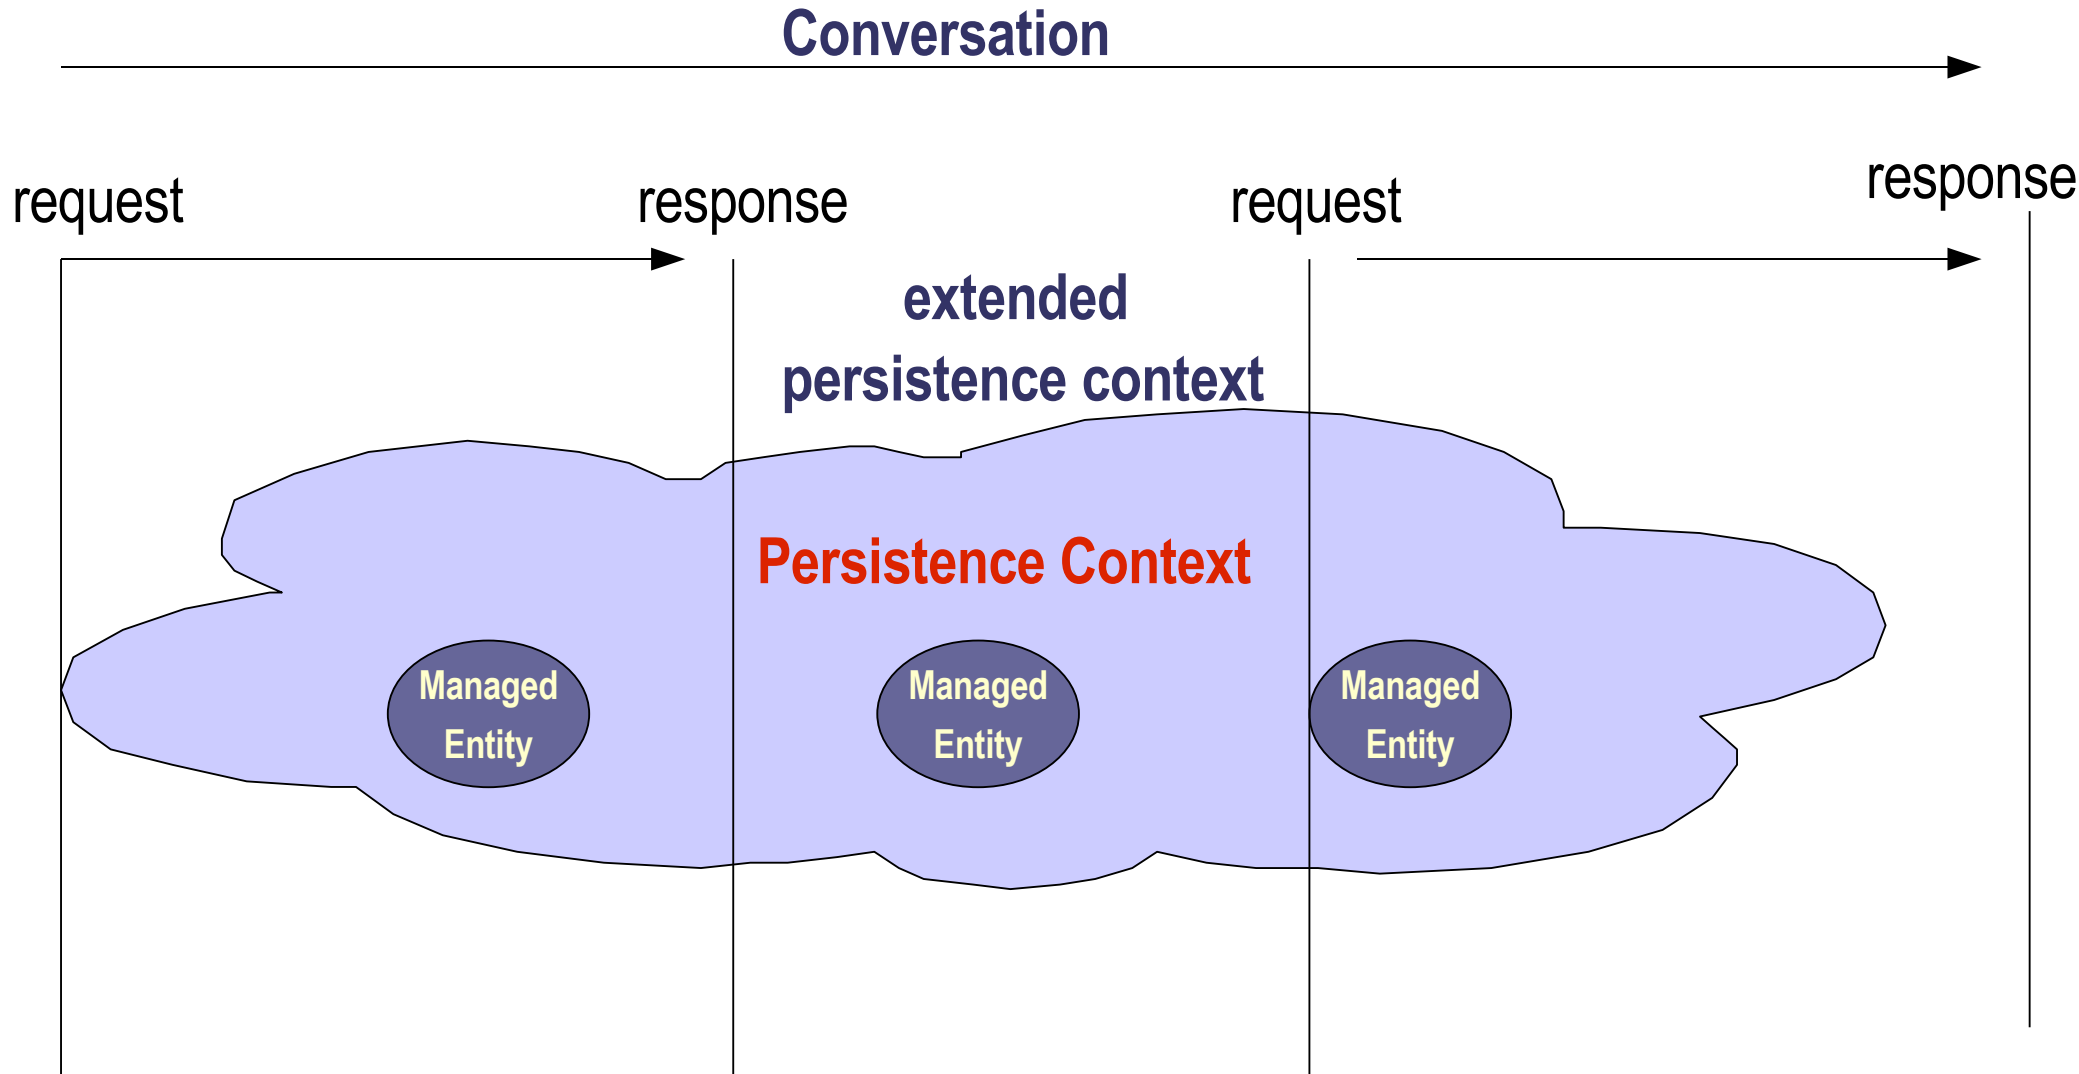


# Types of Persistence Context

- Persistence Context
  - > lifetime maybe transaction-scoped or extended
- Transaction-scoped persistence context
- Extended persistence context
  - > spans **multiple** transactions



# Conversation with Exented Persistence Context



# Extended Persistence Context

```

@Stateful public class OrderMgr {

    //Specify that we want an EXTENDED
    @PersistenceContext(type=PersistenceContextType.EXTENDED)
    EntityManager em;

    //Cached order
    private Order order;

    //create and cache order
    public void createOrder(String itemId) {
        //order remains managed for the lifetime of the bean
        Order order = new Order(cust);
        em.persist(order);
    }

    public void addLineItem(OrderLineItem li) {
        order.lineItems.add(li);
    }

```

Managed entity

Managed entity



# Extended Persistence Context

```

@Stateful public class DeptMgr {
    @PersistenceContext(type=PersistenceContextType.EXTENDED)
    EntityManager em;

    private Department dept;

    @TransactionAttribute(NOT_SUPPORTED)
    public void getDepartment(int deptId) {
        dept = em.find(Department.class, deptId);
    }

    @TransactionAttribute(NOT_SUPPORTED)
    public void addEmployee(int empId) {
        emp = em.find(Employee.class, empId);
        dept.getEmployees().add(emp);
        emp.setDepartment(dept);
    }

    @Remove
    @TransactionAttribute(REQUIRES_NEW)
    public void endUpdate(int deptId) {
        dept = em.find(Department.class, deptId);
    }
}

```



# Persistence Context-Transactional vs. Extended

**@Stateless**

```
public class OrderMgr implements OrderService {
```

```
    @PersistenceContext EntityManager em;
```

```
    public void addLineItem(OrderLineItem li){
```


```
        // First, look up the order.
```

```
        Order order = em.find(Order.class, orderID);
```

```
        order.lineItems.add(li);
```

```
    }
```

look up the order



**@Stateful**

```
public class OrderMgr implements OrderService {
```

```
    @PersistenceContext(type = PersistenceContextType.EXTENDED))
```

```
    EntityManager em;
```

```
    // Order is cached
```

```
    Order order
```

```
    public void addLineItem(OrderLineItem li){
```

```
        // No em.find invoked for the order object
```


```
        order.lineItems.add(li);
```

```
    }
```

Managed entity



No em.find invoked








# Persistence Context Micro Benchmark

- Micro benchmark with lots of lookups
- Persistence context is caching entities





# SEAM Conversations

```

@Name ("shopper")
@Scope (CONVERSATION)
public class BookingManager {


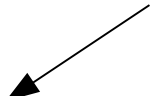

    @In EntityManager entityManager;
    private Booking booking;

    @Begin public void selectHotel(Hotel selectedHotel) {
        hotel = em.merge(selectedHotel);
    }

    @End public void confirm() {
        em.persist(booking);
    }
}
    
```

**SEAM injected**

**SEAM conversation**





# Concurrency and Persistence Context

User 1 transaction


User 2 transaction

Persistence Context 1

Persistence Context 2



same entity



Data source

Object Identity  
only **one** manage entity  
in **PC** represents a row



# Optimistic versus Pessimistic Concurrency

- Optimistic Concurrency
  - > Pros—No database locks held
  - > Cons—Requires a **version attribute** in schema
    - user or app must refresh and retry failed updates
  - > Suitable when application has **few parallel updates**
- Pessimistic Concurrency
  - > Lock the row when data is read in
    - database locks the row upon a select
    - (SELECT . . . FOR UPDATE [NOWAIT])
  - > Pros—Simpler application code
  - > Cons—Database locks
    - limits concurrent access to the data = **scalability**
    - May cause **deadlocks**
    - Not in JPA 1.0 (vendor specific), **supported** in **JPA 2.0**
  - > Suitable when application has **many parallel updates**





# Preventing Parallel Updates

use `@Version` for optimistic locking

```
public class Employee {  
    @ID int id;  
    @Version int version;  
    ...  
}
```

Can be int, Integer,  
short, Short, long,  
Long, Timestamp

Used by persistence manager , Results in following SQL

```
"UPDATE Employee SET ..., version = version + 1  
WHERE id = ? AND version = readVersion"
```

`Version Updated` when transaction commits, merged or  
acquiring a write lock

`OptimisticLockException` if mismatch

Not used by the application!



# Preventing Parallel Updates – 1

```
tx1.begin();
//Joe's employee id is 5
//e1.version == 1
e1 = findPartTimeEmp(5);

//Joe's current rate is $9
e1.raise(2);

tx1.commit();
//e1.version == 2 in db
//Joe's rate is $11
```

Time



```
tx2.begin();
//Joe's employee id is 5
//e1.version == 1
e1 = findPartTimeEmp(5);

//Joe's current rate is $9
if(e1.getRate() < 10)
    e1.raise(5);

//e1.version == 1 in db?
tx2.commit();
//Joe's rate is $14
//OptimisticLockException
```





# Preventing Stale Data JPA 1.0

- Perform **read** or **write locks** on entities
  - Prevents non-repeatable reads in JPA
- entityManager.lock( entity, READ);  
 perform a **version check** on entity **before commit**  
**OptimisticLockException** if mismatch
- entityManager.lock( entity, WRITE);  
 perform a **version check** on entity  
**OptimisticLockException** if mismatch  
**and increment version before commit**



# Preventing Stale Data

```
tx1.begin();
d1 = findDepartment(dId);

//d1's original name is
//"Engrg"
d1.setName("MarketEngrg");

tx1.commit();
```

Time

```
tx2.begin();

e1 = findEmp(eId);
d1 = e1.getDepartment();
em.lock(d1, READ);
if(d1's name is "Engrg")
    e1.raiseByTenPercent();

//Check d1.version in db
tx2.commit();
//e1 gets the raise he does
//not deserve
//Transaction rolls back
```



# Preventing Parallel Updates – 2

Write lock prevents parallel updates

```
tx1.begin();
d1 = findDepartment(dId);

//d1's original name is
//"Engrg"
d1.setName("MarketEngrg");

tx1.commit();
//tx rolls back
```

Time

```
tx2.begin();

e1 = findEmp(eId);
d1 = e1.getDepartment();
em.lock(d1, WRITE);

//version++ for d1
em.flush();
if(d1's name is "Engrg")
    e1.raiseByTenPercent();

tx2.commit();
```



# Bulk Updates

- Update directly against the database, can be Faster But
  - > By pass EntityManager
  - > @Version will not be updated
  - > Entities in PC not updated

```
>tx.begin();
int id = 5; //Joe's employee id is 5
e1 = findPartTimeEmp(id); //Joe's current rate is $9

//Double every employee's salary
em.createQuery(
    "Update Employee set rate = rate * 2").executeUpdate();

//Joe's rate is still $9 in this persistence context
if(e1.getRate() < 10)
    e1.raiseByFiveDollar();

tx.commit();
//Joe's salary will be $14
```



# JPA 2.0 Locks

- JPA1.0 only supported optimistic locking, JPA 2.0 also supports pessimistic locks
- **JPA 2.0** LockMode values :
  - > OPTIMISTIC (= READ)
  - > OPTIMISTIC\_FORCE\_INCREMENT (= WRITE)
  - > PESSIMISTIC
  - > PESSIMISTIC\_FORCE\_INCREMENT
- Multiple places to specify lock

database locks the row  
(SELECT ... FOR UPDATE)



# JPA 2.0 Locking

//Read then lock:

```
Account acct = em.find(Account.class, acctId);
// Decide to withdraw $100 so lock it for update
em.lock(acct, PESSIMISTIC);
int balance = acct.getBalance();
acct.setBalance(balance - 100);
```

Lock after read, risk  
**stale**, could cause  
**OptimisticLock**  
**Exception**

//Read and lock:

```
Account acct = em.find(Account.class,
acctId, PESSIMISTIC);
// Decide to withdraw $100 (already locked)
int balance = acct.getBalance();
acct.setBalance(balance - 100);
```

Locks **longer**,  
could cause  
bottlenecks,  
deadlock





# JPA 2.0 Locking

Trade-offs:

- lock earlier : **risk bad scalability, deadlock**
- Lock later : risk **stale** data for update, get optimistic lock exception

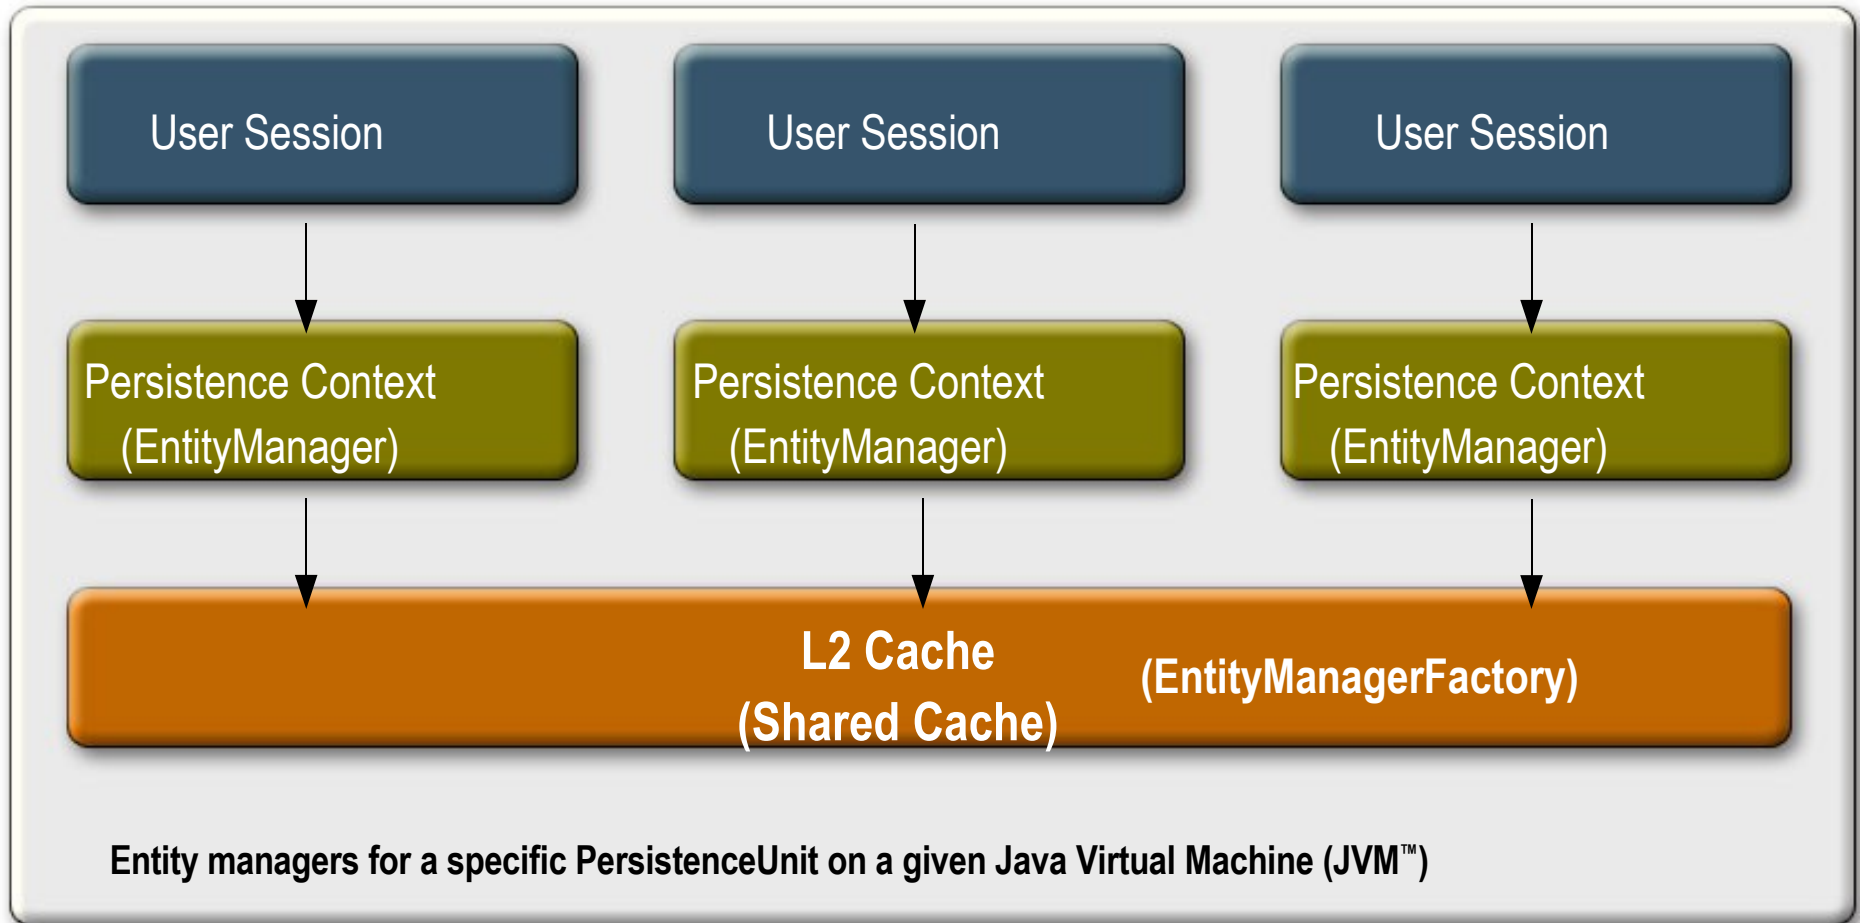
```
// read then lock and refresh
Account acct = em.find(Account.class, acctId);
// Decide to withdraw $100 - lock and refresh
em.refresh(acct, PESSIMISTIC);
int balance = acct.getBalance();
acct.setBalance(balance - 100);
```

“right” approach depends on requirements



# L2 cache shared across transactions and users

Putting it all together





## Second-level Cache

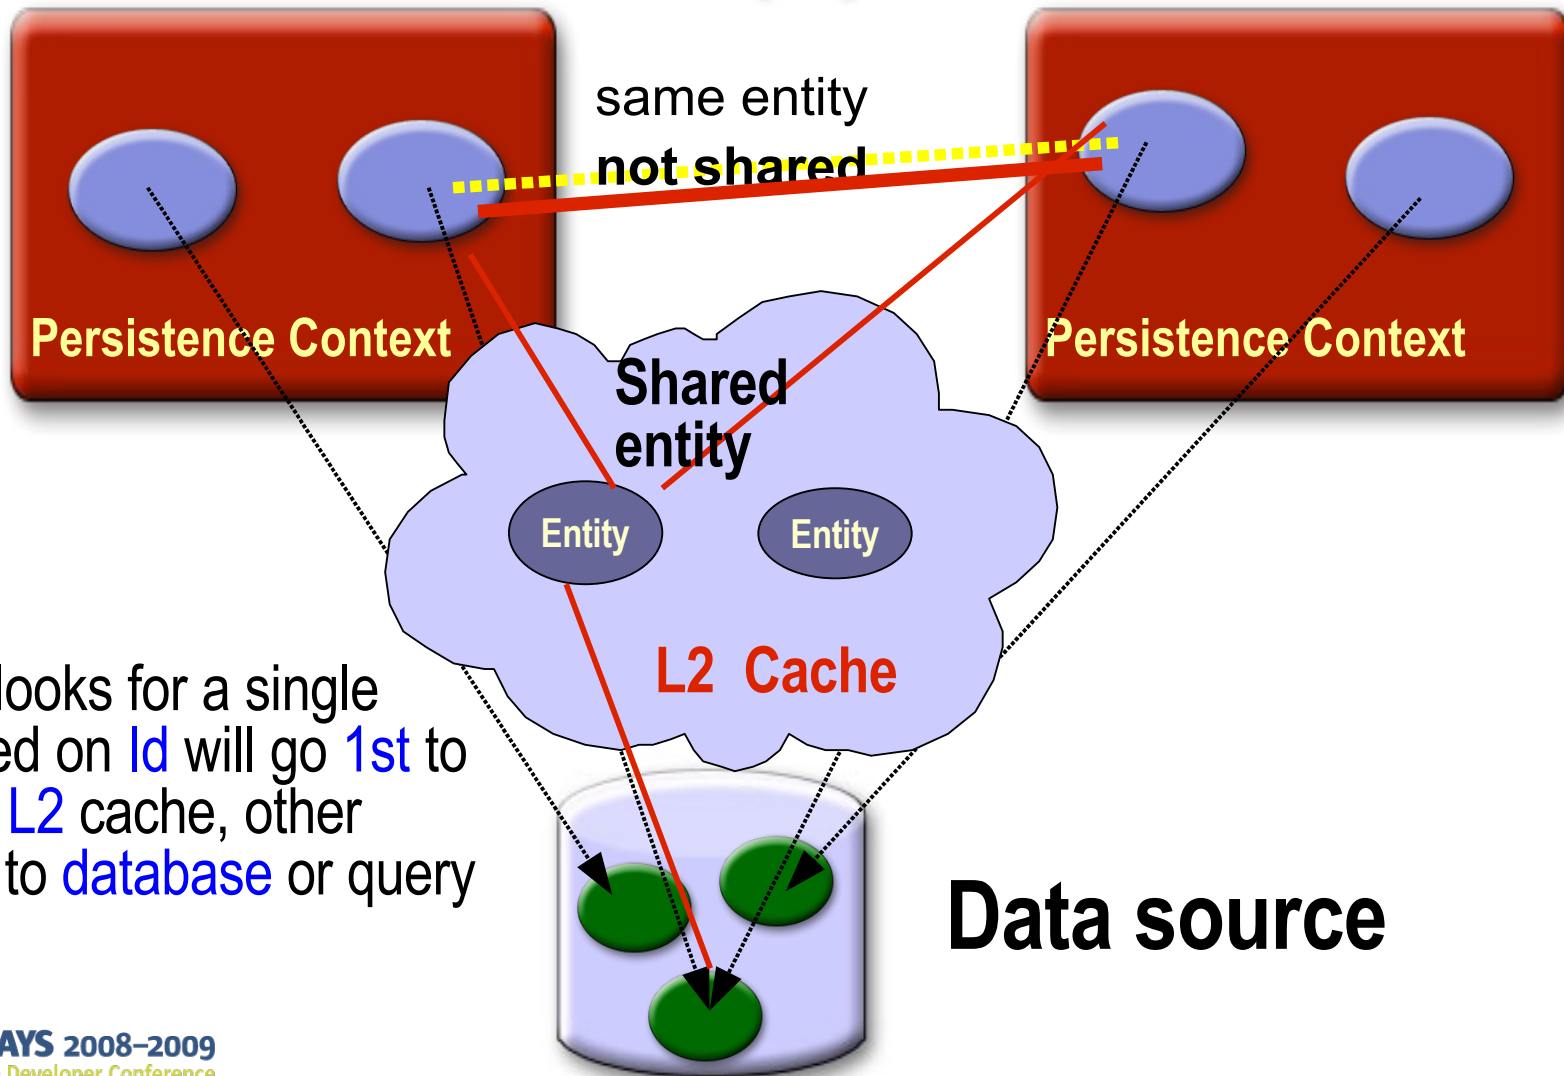
- L2 Cache shares entity state across various persistence contexts
  - > If caching is enabled, entities **not found** in **persistence context**, will be **loaded from L2 cache**, if found
- Best for **read-mostly** classes
- L2 Cache is Vendor specific
  - > Java Persistence API 1.0 does not specify L2 support
  - > Java Persistence API 2.0 has basic cache operations
  - > Most persistence providers-- Hibernate, EclipseLink, OpenJPA ... provide second-level cache(s)



# L2 Cache

User transaction 1

User transaction 2



query that looks for a single object based on **id** will go **1st** to **PC** then to **L2** cache, other queries go to **database** or query cache



## L2 Caching

- **Pros:**
  - > avoids database access for already loaded entities
    - faster for reading frequently accessed unmodified entities
- **Cons**
  - > memory consumption for large amount of objects
  - > Stale data for updated objects
  - > Concurrency for write (optimistic lock exception, or pessimistic lock)
    - Bad scalability for frequent or concurrently updated entities



## L2 Caching

- Configure L2 caching for entities that are
  - > read often
  - > modified infrequently
  - > Not critical if stale
- protect any data that can be concurrently modified with a locking strategy
  - > Must handle optimistic lock failures on flush/commit
  - > configure expiration, refresh policy to minimize lock failures
- Configure Query cache
  - > Useful for queries that are run frequently with the same parameters, for not modified tables





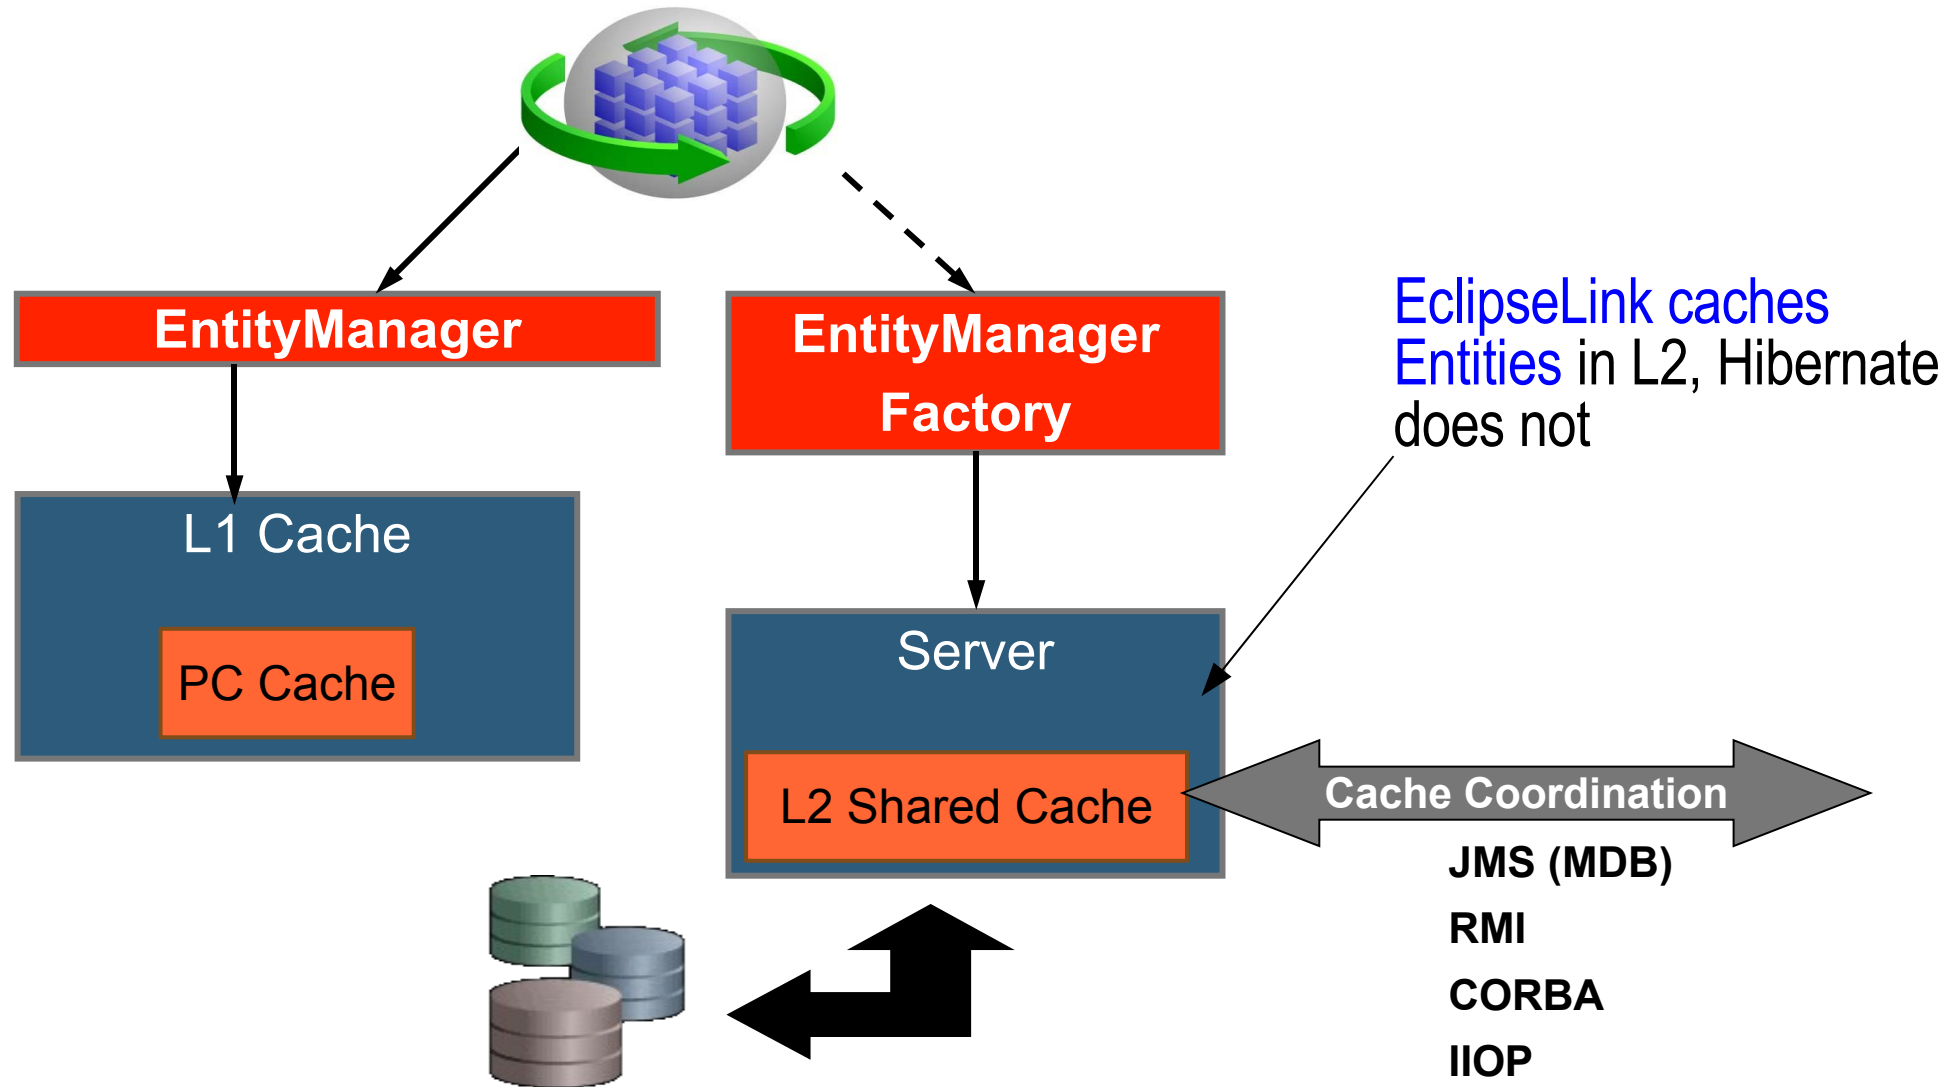
## JPA 2.0 Shared Cache API

- entity cache shared across persistence unit
  - > Accessible from `EntityManagerFactory`
- Supports only very **basic** cache operations
  - > Can be extended by vendors

```
public class Cache {
    //checks if object is in IdentityMap
    public boolean contains(Class class, Object pk);
    // invalidates object in the IdentityMap
    public void evict(Class class, Object pk);
    public void evict(Class class); // invalidates the class in the IdentityMap.
    public void evictAll(); // Invalidates all classes in the IdentityMap
}
```



# EclipseLink Caching Architecture





## EclipseLink Extensions - L2 Caching

- **Default:** Entities read are L2 **cached**
- Cache Configuration by Entity type or Persistence Unit
  - > **You can disable L2 cache**
- Configuration Parameters
  - > Cache isolation, type, size, expiration, coordination, invalidation, refreshing
  - > Coordination (cluster-messaging)
    - Messaging: JMS, RMI, RMI-IIOP, ...
    - Mode: SYNC, SYNC+NEW, INVALIDATE, NONE



# EclipseLink Mapping Extensions

```
@Entity
@Table(name="EMPLOYEE")
@Cache (
    type=CacheType.WEAK,
    isolated=false,
    expiry=600000,
    alwaysRefresh=true,
    disableHits=true,
    coordinationType=INVALIDATE_CHANGED_OBJECTS
)
public class Employee implements Serializable {
    ...
}
```

## Type=

**Full:** objects never flushed unless deleted or evicted

**weak:** object will be garbage collected if not referenced

**=true**

disables L2 cache

## @Cache

- type, size, isolated, expiry, refresh, cache usage, coordination
- Cache usage and refresh query hints



# Hibernate L2 Cache

- Hibernate L2 cache is **not configured by default**
- Hibernate L2 does not cache Entities. Hibernate caches Id and state
- Hibernate L2 cache is **pluggable**
  - > EHCache, OSCache, SwarmCacheProvider (**JVM**)
  - > JBoss TreeCache **Cluster**
  - > Can plug in others like Terracotta

## Cache Concurrency Strategy

Cache	Type	Read-only	Read-write	Transactional
<b>EHCache</b>	memory, disk	Yes	Yes	
<b>OSCache</b>	memory, disk	Yes	Yes	
<b>SwarmCache</b>	clustered	Yes		
<b>JBoss Cache</b>	clustered	Yes		Yes



# Hibernate L2 Cache

not configured by default

<!-- optional configuration file parameter -->

net.sf.ehcache.configurationResourceName=/name\_of\_configuration\_resource

@Entity

@Cache(usage =  
CacheConcurrencyStrategy.NONSTRICT\_READ\_WRITE)

```
public Class Country {
    private String name;
    ...
}
```

Cache Concurrency Strategy  
must be supported by cache provider



# OpenJPA L2 Caching

- OpenJPA L2 caches object data and JPQL query results
- Updated when data is loaded from database and after changes are successfully committed
- For cluster caches are notified when changes are made
- Can plug in implementations, such as Tangosol's Coherence product
- several cache eviction strategies:
  - > Time-to-live settings for cache instances
 

```
@Entity @DataCache(timeout=5000)
public class Person { ... }
```





# Agenda


- > Entity Manager
- > Persistence Context
- > **Entities**
- > Schema and Queries
- > Transaction



# Maintaining Relationship

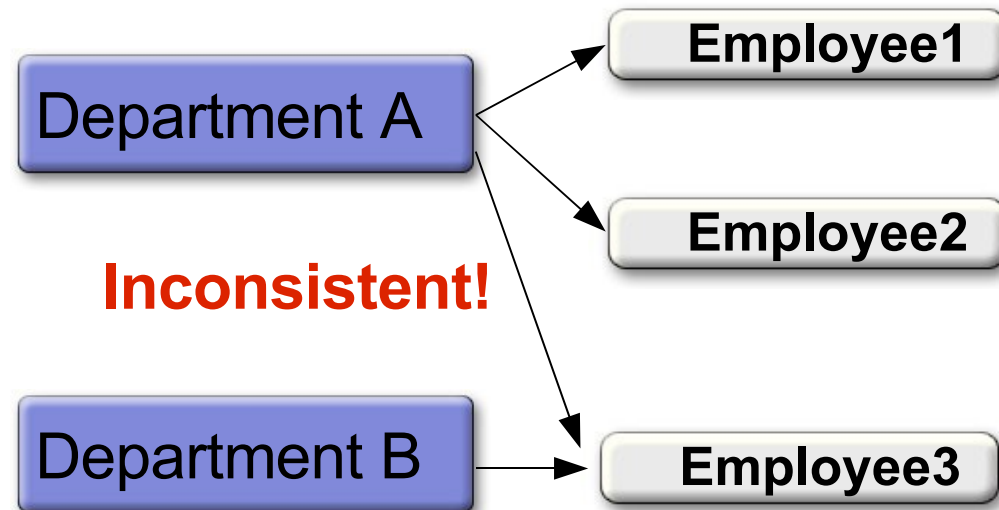
- Application bears the responsibility of maintaining relationship between objects

## Before



```
deptA.getEmployees().add(e3);
```

## After





## Example – Domain Model

```
@Entity public class Employee {
    @Id private int id;
    private String firstName;
    private String lastName;
    @ManyToOne(fetch=LAZY)
    private Department dept;
    ...
}

@Entity public class Department {
    @Id private int id;
    private String name;
    @OneToMany(mappedBy = "dept", fetch=LAZY)
    private Collection<Employee> emps = new ...;
    ...
}
```



## Example – Managing Relationship

# INCORRECT

```
public int addNewEmployee(...) {  
    Employee e = new Employee(...);  
    Department d = new Department(1, ...);  
  
    e.setDepartment(d);  
    //Reverse relationship is not set  
    em.persist(e);  
    em.persist(d);  
  
    return d.getEmployees().size();  
}
```



## Example – Managing Relationship

**CORRECT**

```
public int addNewEmployee(...) {
    Employee e = new Employee(...);
    Department d = new Department(1, ...);

    e.setDepartment(d);
    d.getEmployees().add(e);
    em.persist(e);
    em.persist(d);

    return d.getEmployees().size();
}
```



# Navigating Relationships

## Data fetching strategy

- EAGER – immediate
- LAZY – loaded only when needed
- LAZY is good for large objects and/or with relationships with deep hierarchies



# Lazy loading and JPA

```
@Entity public class Department {
    @Id private int id;
    @OneToMany(mappedBy = "dept")
    private Collection<Employee> emps ;
    ...
}
```

- Default FetchType is LAZY for 1:m and m:n relationships
  - > benefits large objects and relationships with deep hierarchies
- However for use cases where data is needed can cause **n+1** selects
- LAZY – N + 1 problem:

```
SELECT d.id, ... FROM Department d // 1 time
SELECT e.id, ... FROM Employee e
WHERE e.deptId = ? // N times
```





# Lazy loading and JPA

```
@Entity public class Department {
    @Id private int id;
    @OneToMany(mappedBy = "dept", fetch=EAGER)
    private Collection<Employee> employees ;
    ...
}
```

- Relationship can be Loaded Eagerly Can cause Cartesian product
  - > But if you have several **related relationships**, could **load too much** !

OR

- Temporarily override the LAZY fetch type, use **Join Fetch** in a **query**:

```
@NamedQueries({ @NamedQuery(name="getItEarly",
    query="SELECT d FROM Department d JOIN FETCH d.employees") })
```

```
public class Department{
    .....
}
```



# Lazy loading and JPA

- Capture generated SQL
  - > persistence.xml file: `<property name="toplink.logging.level" value="FINE">`
- Test run use cases and examine the SQL statements
  - > **optimise the number** of SQL statements executed!
  - > **only retrieve the data** your application **needs!**
- **Lazy** load **large** (eg BLOB) attributes and relationships that are not used often
- Override to Eagerly load in use cases where needed



# Navigating Relationships

Detached  
Entity

- Accessing a **LAZY relationship** from a **detached** entity
  - > If not loaded , Causes an exception
- Solutions:
  - > Use JOIN FETCH
  - > Or Set Fetch type to EAGER
  - > Or Access the collection before entity is detached:

```
d.getEmployees().size();
```





# Navigating Relationships

## Data fetching strategy

- Cascade specifies operations on relationships
  - > ALL, PERSIST, MERGE, REMOVE, REFRESH
  - > The default is do **nothing**
- Avoid **MERGE, ALL** with deep hierarchies
  - > If want to do it, limit the scope



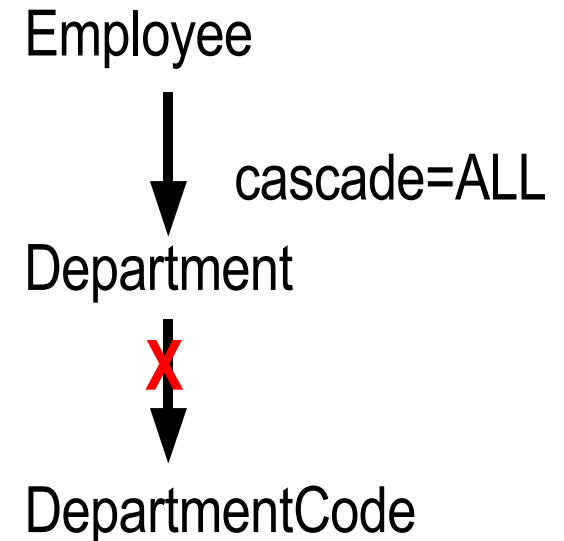
# Using Cascade

```

@Entity public class Employee {
    @Id private int id;
    private String firstName;
    private String lastName;
    @ManyToOne(cascade=ALL, fetch=LAZY)
    private Department dept;
    ...
}

@Entity public class Department {
    @Id private int id;
    private String name;
    @OneToMany(mappedBy = "dept"
                cascade=ALL, fetch=LAZY)
    private Collection<Employee> emps = new ...;
    @OneToMany
    private Collection<DepartmentCode> codes;
    ...
}

```





# Agenda

- > Entity Manager
- > Persistence Context
- > Entities
- > **Schema and Queries**
- > Transaction

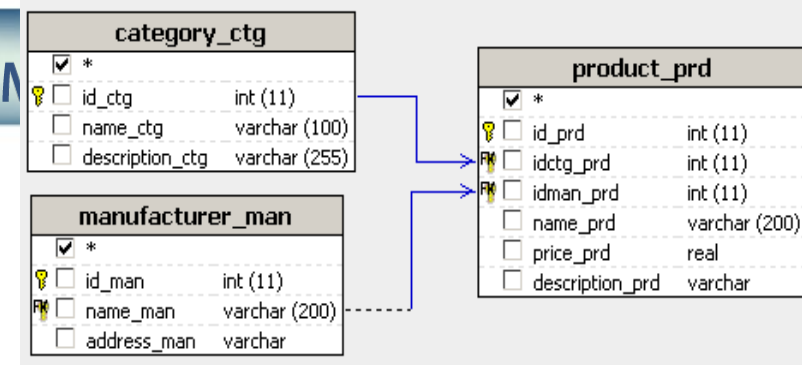


# Database design Basic foundation of performance

- **Smaller tables use less disk, less memory, can give better performance**
  - > **Use as small data types as possible**
  - > **use as small primary key as possible**
  - > **Vertical Partition:**
    - **split large, infrequently used columns into a separate one-to-one table**
- **Use good indexing**
  - > **Indexes Speed up Querys**
  - > **Indexes slow down Updates**
  - > **Index columns frequently used in Query Where clause**



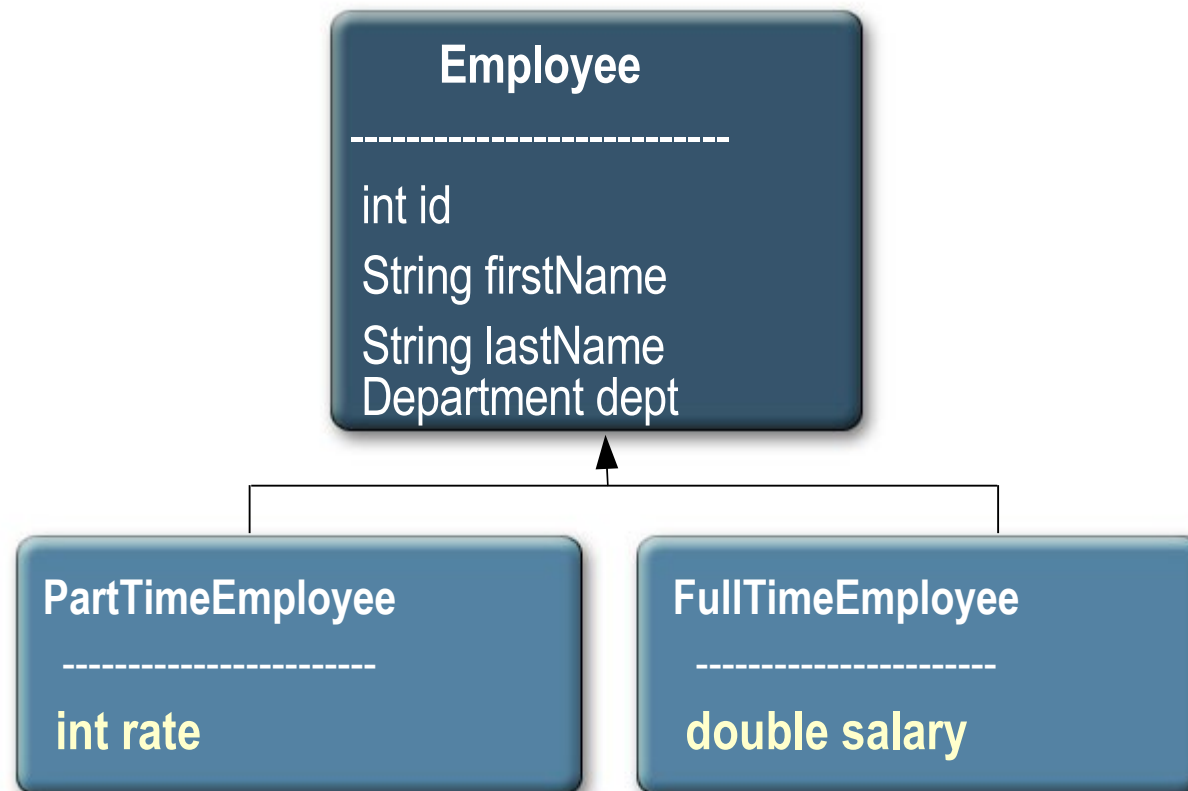
# Normalization



- Normalization Eliminates redundant data
  - > updates are usually **faster**.
    - there's **less data to change**.
- However Normalized database **causes joins** for queries
  - > Queries maybe slower
  - > Normalize then maybe De-normalize frequently read columns and cache them



# Mapping Inheritance Hierarchies





# Single Table Per Class

## Benefits

@Inheritance(strategy=SINGLE\_TABLE)

- Simple
- **No joins** required
  - > can be **fast** for **Queries**

## Drawbacks

- **Not normalized**
  - > **Wastes space**
- Requires columns corresponding to subclasses' state to be null
- Table can have too many columns
  - > **Larger tables**= more data, can have negative affects on performance

EMPLOYEE	
-----	
ID	Int PK,
FIRSTNAME	varchar(255),
LASTNAME	varchar(255),
DEPT_ID	int FK,
<b>RATE</b>	<b>int NULL,</b>
<b>SALARY</b>	<b>double NULL,</b>
<b>DISCRIM</b>	<b>varchar(30)</b>



# Joined Subclass

@Inheritance(strategy=JOINED)

## Benefits

- **Normalized** database
  - > Better for storage
- Database view same as domain model
- Easy to evolve domain model

## Drawbacks

- Queries cause **joins**
  - > **Slower queries**
  - > Poor performance for deep hierarchies, polymorphic queries and relationships

EMPLOYEE	
ID	int PK,
FIRSTNAME	varchar(255),
LASTNAME	varchar(255),
DEPT_ID	int FK,

PARTTIMEEMPLOYEE	
ID	int PK FK,
<b>RATE</b>	<b>int NULL</b>

FULLTIMEEMPLOYEE	
ID	int PK FK,
<b>SALARY</b>	<b>double NULL</b>



# Table Per Class

@Inheritance(strategy=TABLE\_PER\_CLASS)

## Benefits

- No need for joins to read entities of same type
  - > Faster reads

## Drawbacks

- Not normalized
  - > Wastes space
- Polymorphic queries cause **union** (all employees)
  - > Poor performance
- This strategy is not mandatory

PARTTIMEEMPLOYEE	
ID	int PK,
FIRSTNAME	varchar(255),
LASTNAME	varchar(255),
DEPT_ID	int FK,
RATE	int NULL


FULLTIMEEMPLOYEE	
ID	int PK,
FIRSTNAME	varchar(255),
LASTNAME	varchar(255),
DEPT_ID	int FK,
SALARY	double NULL



# vertical partitioning

```
CREATE TABLE Customer (
  user_id INT NOT NULL AUTO_INCREMENT
, email VARCHAR(80) NOT NULL
, display_name VARCHAR(50) NOT NULL
, password CHAR(41) NOT NULL
, first_name VARCHAR(25) NOT NULL
, last_name VARCHAR(25) NOT NULL
, address VARCHAR(80) NOT NULL
, city VARCHAR(30) NOT NULL
, province CHAR(2) NOT NULL
, postcode CHAR(7) NOT NULL
, interests TEXT NULL
, bio TEXT NULL
, signature TEXT NULL
, skills TEXT NULL
, PRIMARY KEY (user_id)
, UNIQUE INDEX (email)
) ENGINE=InnoDB;
```

Frequently  
referenced



Less Frequently  
referenced,  
TEXT data

```
CREATE TABLE Customer(
  user_id INT NOT NULL AUTO_INCREMENT
, email VARCHAR(80) NOT NULL
, display_name VARCHAR(50) NOT NULL
, password CHAR(41) NOT NULL
, PRIMARY KEY (user_id)
, UNIQUE INDEX (email)
) ENGINE=InnoDB;
```

```
CREATE TABLE CustomerInfo (
  user_id INT NOT NULL
, first_name VARCHAR(25) NOT NULL
, last_name VARCHAR(25) NOT NULL
, address VARCHAR(80) NOT NULL
, city VARCHAR(30) NOT NULL
, province CHAR(2) NOT NULL
, postcode CHAR(7) NOT NULL
, interests TEXT NULL
, bio TEXT NULL
, signature TEXT NULL
, skills TEXT NULL
, PRIMARY KEY (user_id)
, FULLTEXT KEY (interests, skills)
) ENGINE=MyISAM;
```

- limit number of columns per table
- split large, infrequently used columns into a separate table



# Vertical partitioning

**@Entity**

```
public class Customer {
    int userid;
    String email;
    String password;
    @OneToOne(fetch=LAZY)
    CustomerInfo info;
}
```

**@Entity**

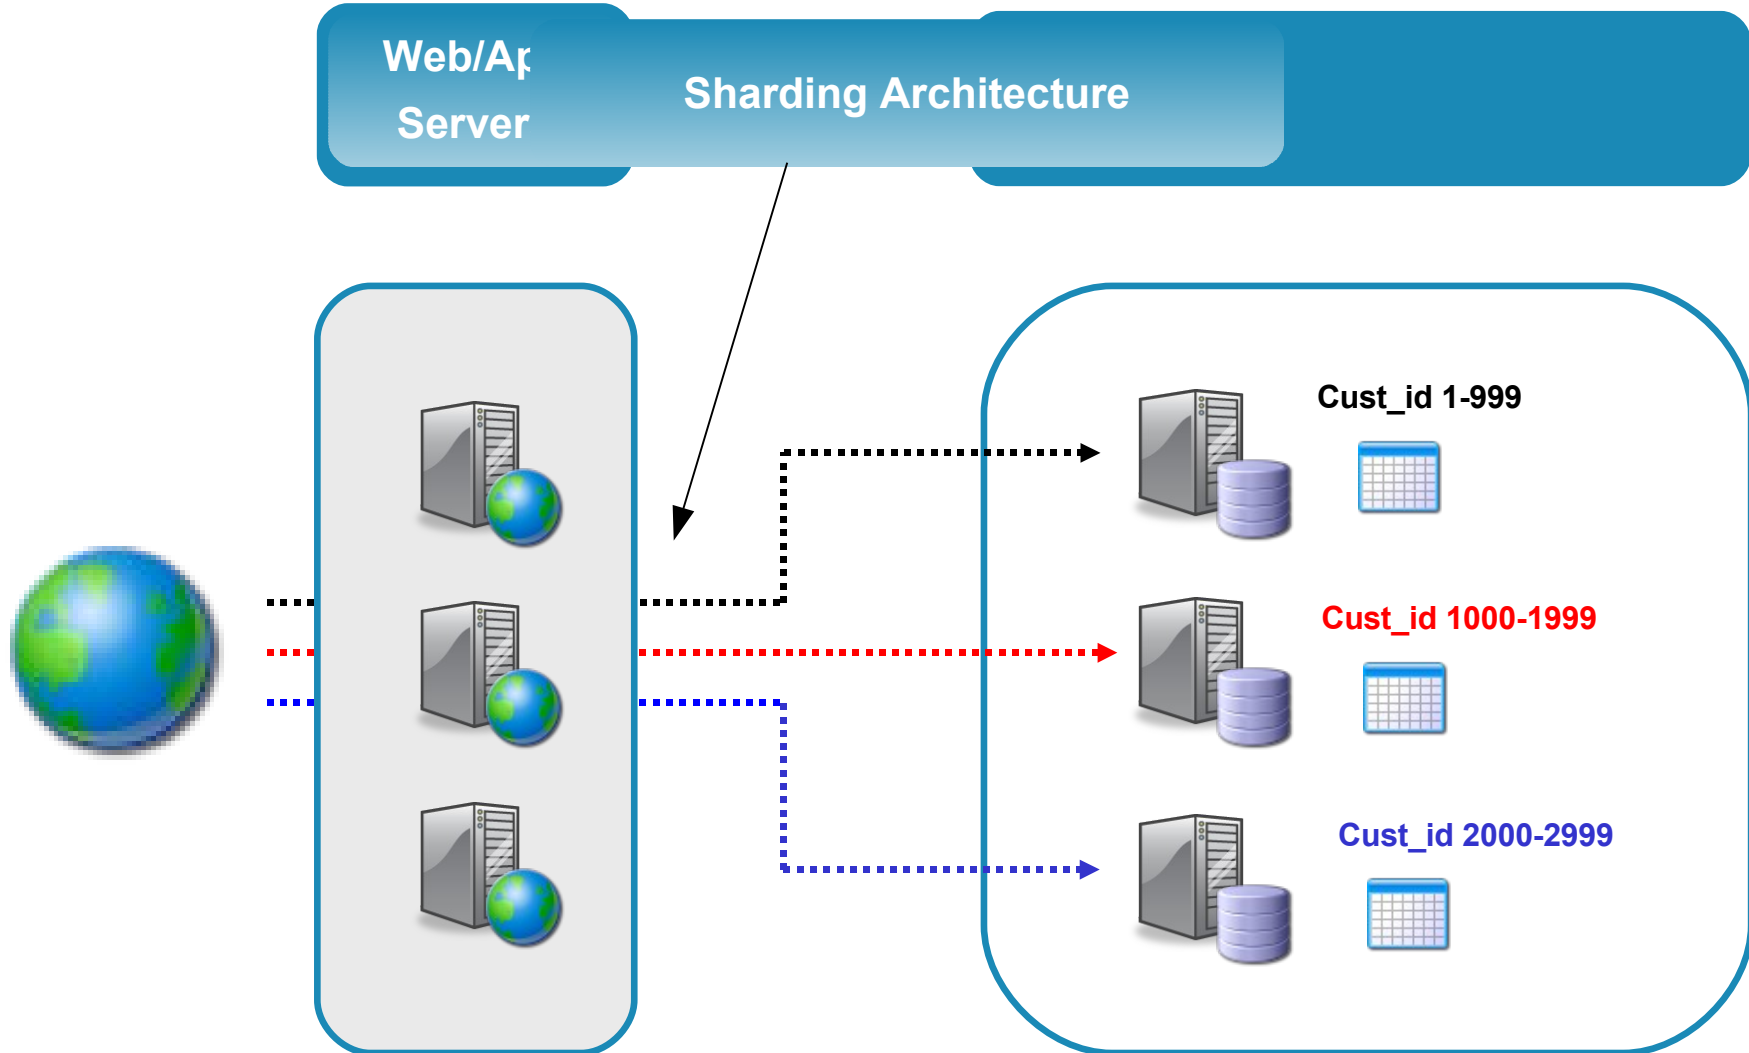
```
public class CustomerInfo {
    int userid;
    String name;
    String interests;
    @OneToOne(mappedBy=
        "CustomerInfo")
    Customer customer;
}
```

- split large, infrequently used columns into a separate table





# Scalability: Sharding - Application Partitioning



> Sharding = Horizontal partitioning

- Split table by rows into partitions



# Know what SQL is executed

- Capture generated SQL:  
persistence.xml file:  
<property name="toplink.logging.level" value="FINE">
- Find and fix problem SQL:
  - > Watch for **slow** Queries
    - use the MySQL **slow query log** and use Explain
      - Can reveal problems such as a **missing Indexes**
  - > Watch for Queries that execute **too often** to load needed data
  - > Watch for loading more data than needed



# MySQL Query Analyser

Monitor

Advisors

Events

Graphs

Query Analysis

Replication

Settings

What's New

troubleshooting

configure query analysis

All Servers

Browse Queries

[ 1 to 20 of 47 ]

1 2 3 next » last »

Search Type

Query Search

Database

Time Display

Hours

Minutes

View

Query Type

Limit

Contains

Interval

00

30

Group

All

20

filter

reset

Query	Database	Exec Count	Exec Time (hh:mm:ss.ms)			Rows			Bytes			First Seen	
			Total	Max	Avg	Total	Max	Avg	Total	Max	Avg		
<a href="#">SELECT COUNT( message_i... process_type , fmtdate (1)</a>	intranet_mcslp	6	12.838	11.123	2.140	776	408	129	24.21 KB	12.8 KB	4.04 KB	10:35:01 AM	<a href="#">alias</a>
<a href="#">SELECT inhost , path , ...GROUP BY fmtdate , path (1)</a>	intranet_mcslp	2	0.906	0.686	0.453	4,038	2,526	2,019	230.89 KB	142.05 KB	115.44 KB	10:34:01 AM	<a href="#">alias</a>
<a href="#">SELECT process_mode , p...ess_mode , process_type (1)</a>	intranet_mcslp	6	1.674	0.333	0.279	36	7	6	808 B	165 B	134.67 B	10:37:01 AM	<a href="#">alias</a>
<a href="#">SELECT media_photo . ph...RDER BY RAND( ) LIMIT ? (1)</a>	intranet_mcslp	1	0.259	0.259	0.259	30	30	30	282 B	282 B	282 B	10:37:01 AM	<a href="#">alias</a>
<a href="#">SELECT process_mode , p...ess_mode , process_type (1)</a>	intranet_mcslp	1	0.241	0.241	0.241	8	8	8	189 B	189 B	189 B	10:33:01 AM	<a href="#">alias</a>
<a href="#">SELECT data , COUNT( DI...P BY data ORDER BY data (1)</a>	intranet_mcslp	1	0.217	0.217	0.217	195	195	195	3.06 KB	3.06 KB	3.06 KB	10:39:01 AM	<a href="#">alias</a>
<a href="#">SELECT DISTINCT( media... , photoid DESC LIMIT ? (1)</a>	intranet_mcslp	25	1.781	0.209	0.071	250	10	10	2.44 KB	100 B	100 B	10:38:01 AM	<a href="#">alias</a>
<a href="#">SELECT DISTINCT( album ...RDER BY RAND( ) LIMIT ? (1)</a>	intranet_mcslp	1	0.193	0.193	0.193	10	10	10	594 B	594 B	594 B	10:37:01 AM	<a href="#">alias</a>
<a href="#">INSERT INTO currencies ...ALUES ( ? , ? , ? , ? ) (1)</a>	intranet_mcslp	5	0.109	0.105	0.022	5	1	1	0 B	0 B	0 B	10:31:01 AM	<a href="#">alias</a>
<a href="#">SELECT COUNT( DISTINCT( ... meta WHERE type = ? (1)</a>	intranet_mcslp	1	0.095	0.095	0.095	1	1	1	8 B	8 B	8 B	10:37:01 AM	<a href="#">alias</a>

## Find and fix problem SQL:

- how long a query took
- results of EXPLAIN statements
- Historical and real-time analysis
  - > **query execution counts**, run time

Its **not just slow running** queries that are a problem, Sometimes its **SQL** that **executes a lot** that kills your system



# Agenda

- > Entities
- > Entity Manager
- > Persistence Context
- > **Queries**
- > Transaction



# Query Types – 1

- Named query
  - > Like `findByXXXX()` from `EntityBeans`
  - > Compiled by persistence engine
  - > Created either with `@NamedQuery` or externalized in `orm.xml`

- Dynamic query

- > Created dynamically by passing a query string to `EntityManager`

```
Query query = em.createQuery("select ...");
```

- > Beware of **SQL injection**, better to use with named parameters

**NOT GOOD**

```
q = em.createQuery("select e from Employee e WHERE "
    + "e.empId LIKE '" + id + "'");
```

**GOOD**

```
q = em.createQuery("select e from Employee e WHERE "
    + "e.empId LIKE ':id'");
```

```
q.setParameter("id", id);
```

## Query Types – 2

- Native query
  - > Leverage the native database querying facilities
  - > Not portable – ties queries to database



# Flush Mode

- Controls whether the state of managed entities are synchronized before a query
- Types of flush mode
  - > AUTO – immediate, default
  - > COMMIT – flush only when a transaction commits
  - > NEVER – need to invoke `EntityManager.flush()` to flush

```
//Assume JTA transaction
Order order = em.find(Order.class, orderNumber);
em.persist(order);
Query q = em.createNamedQuery("findAllOrders");
q.setParameter("id", orderNumber);
q.setFlushMode(FlushModeType.COMMIT);
//Newly added order will NOT visible
List list = q.getResultList();
```





# Agenda

- > Entities
- > Entity Manager
- > Persistence Context
- > Queries
- > **Transaction**



# Transactions

- Do not perform expensive and unnecessary operations that are not part of a transaction
  - > Hurt performance
  - > Eg. logging – disk write are expensive, resource contention on log
- Do not use transaction when browsing data
  - > **@TransactionAttribute(NOT\_SUPPORTED)**

# THANK YOU

Carol McDonald  
Java Architect

**SUN TECH DAYS 2008–2009**  
A Worldwide Developer Conference